



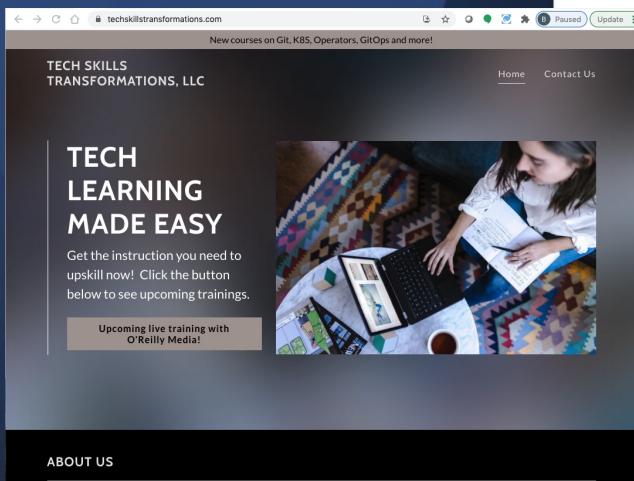
# Prerequisites

- For this workshop, you can use either the preconfigured VM environment or a user-supplied environment.
- For details on the VM option, see the file [\*\*k8s-dev-setup-vm-environment-option.pdf\*\*](#) in [github.com/skilldocs/k8s-dev](https://github.com/skilldocs/k8s-dev)
  - For this option, you will need VirtualBox and the .ova file
    - » <https://www.dropbox.com/s/d8bjycx07v9is4i/k8s-dev.ova?dl=0>
    - OR -
    - » <https://bclconf.s3.us-west-2.amazonaws.com/k8s-dev.ova>
- For details on the user-supplied option, see the file [\*\*k8s-dev-setup-user-supplied-environment-option.pdf\*\*](#) in [github.com/skilldocs/k8s-dev](https://github.com/skilldocs/k8s-dev)
- You will also need the labs doc for the workshop from <https://github.com/skilldocs/k8s-dev/blob/main/k8s-dev.pdf>

# Kubernetes for Developers Deep Dive

Brent Laster  
Tech Skills Transformations

# About me



- Founder, Tech Skills Transformations LLC
- R&D DevOps Director
- Global trainer – training (Git, Jenkins, Gradle, CI/CD, pipelines, Kubernetes, Helm, ArgoCD, operators)
- Author -
  - OpenSource.com
  - Professional Git book
  - Jenkins 2 – Up and Running book
  - Continuous Integration vs. Continuous Delivery vs. Continuous Deployment mini-book on O'Reilly
- <https://www.linkedin.com/in/brentlaster>
- @BrentCLaster
- GitHub: brentlaster



# Professional Git Book

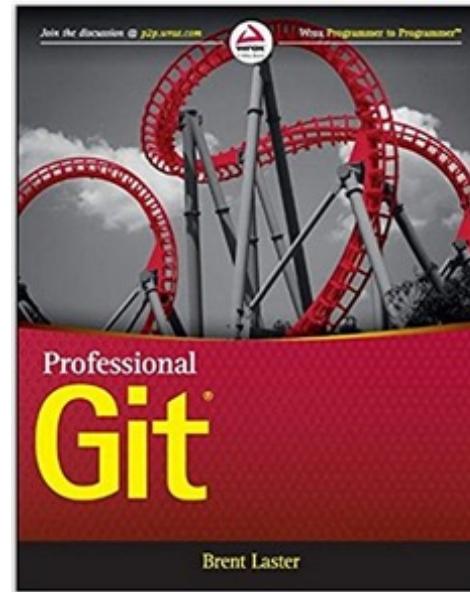
- Extensive Git reference, explanations,
- and examples
- First part for non-technical
- Beginner and advanced reference
- Hands-on labs

## Professional Git 1st Edition

by [Brent Laster](#) (Author)

7 customer reviews

[Look inside](#)



© 2023 Brent C. Laster &

Tech Skills Transformations LLC



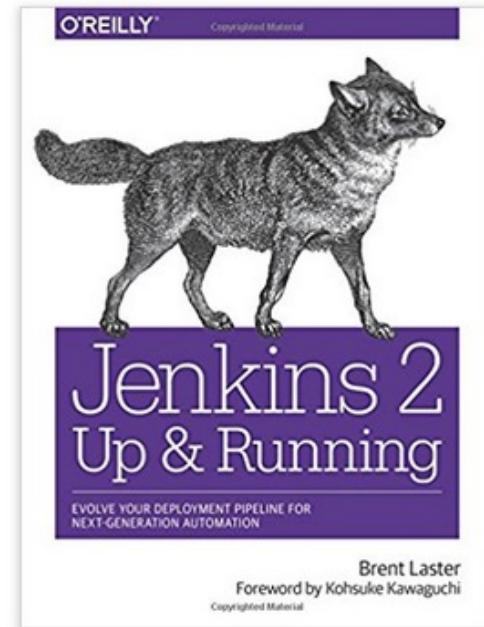
# Jenkins 2 Book

- Jenkins 2 – Up and Running
- “It’s an ideal book for those who are new to CI/CD, as well as those who have been using Jenkins for many years. This book will help you discover and rediscover Jenkins.” *By Kohsuke Kawaguchi, Creator of Jenkins*

★★★★★ 5 customer reviews

#1 New Release in Java Programming

[Look inside](#) ↴





# Learning GitHub Actions - Early Release

7

- Early release - Chapters 1-6 on learning.oreilly.com

**O'REILLY®** TEAMS ▾ INDIVIDUALS FEATURES ▾ BLOG CONTENT SPONSORSHIP SIGN IN TRY NOW >

See everything available through the O'Reilly learning platform and start a free trial. Explore now.

**Search**

**Learning GitHub Actions**  
Automation and Integration of CI/CD with GitHub  
by [Brent Laster](#)  
Released September 2023  
Publisher(s): O'Reilly Media, Inc.  
ISBN: 9781098131074

Read it now on the O'Reilly learning platform with a 10-day free trial.

O'Reilly members get unlimited access to books, live events, courses curated by job role, and more from O'Reilly and nearly 200 top publishers.

**START YOUR FREE TRIAL >**



# O'Reilly Training

April 26 & 27, 2021

## Git Fundamentals

Join Brent Laster to gain the knowledge and skills you need to leverage Git to greatly simplify and speed up managing all of the changes in your source code. Once you ...

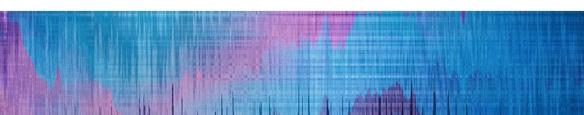
### LIVE ONLINE TRAINING Next Level Git - Master your content

Use powerful tools in Git to simplify merges, rewrite history, and perform automatic updates

Topic: Software Development



BRENT LASTER



### LIVE ONLINE TRAINING Building a deployment pipeline with Jenkins 2

Manage continuous integration and continuous delivery to release software

Topic: System Administration



BRENT LASTER



LIVE ONLINE TRAINING

Containers A-Z

An overview of containers, Docker, Kubernetes, Istio, Helm, Kubernetes Operators, and GitOps

Topic: System Administration



BRENT LASTER



June 3, 10, 17 & 24, 2021

## Git in 4 Weeks

Join author, trainer, and DevOps director Brent Laster to learn how Git works—and discover all the things you can do with it. Over four sessions, Brent walks you through everything you ...

### LIVE ONLINE TRAINING Next Level Git - Master your workflow

Use Git to find problems, simplify working with multiple branches and repositories, and customize behavior with hooks

Topic: Software Development



BRENT LASTER



LIVE ONLINE TRAINING

## Git Troubleshooting

How to solve practically any problem that comes your way

Topic: Software Development



BRENT LASTER



### LIVE ONLINE TRAINING

## Getting started with continuous delivery (CD)

Move beyond CI to build, manage, and deploy a working pipeline

Topic: System Administration



BRENT LASTER



### LIVE ONLINE TRAINING

## Helm Fundamentals

Deploying, upgrading, and rolling back applications in Kubernetes

Topic: System Administration



BRENT LASTER



June 28, 2021

## Troubleshooting Kubernetes

In this 3-hour course, global trainer, author, and DevOps director Brent Laster will show you how to respond to the most common problem situations you may encounter with Kubernetes. You'll learn ...



May 24, 2021

## Continuous Delivery in Kubernetes with ArgoCD

Join expert Brent Laster to explore GitOps and learn how to use Argo CD to implement GitOps in your Kubernetes deployments. APAC time friendly - You're a Kubernetes admin who wants ...

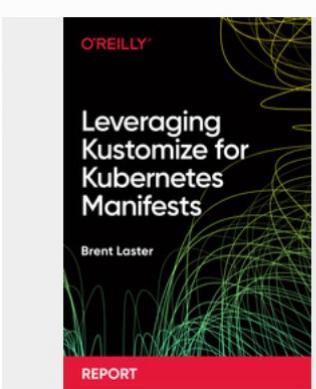


May 17, 2021

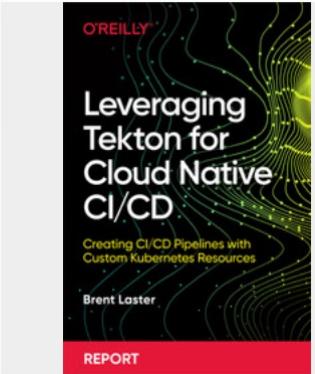
## Building a Kubernetes Operator

Join expert Brent Laster to learn how the Operator pattern helps address these kinds of situations by allowing you to create custom controllers that extend the functionality of the Kubernetes API ...

# Other References



[Leveraging Kustomize for Kubernetes Manifests](#)



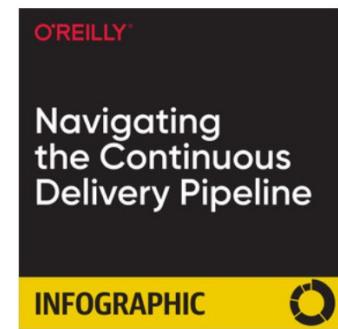
[Leveraging Tekton for Cloud Native CI/CD](#)



[Continuous Integration vs. Continuous Delivery vs. Continuous Deployment, 2nd Edition](#)



[Continuous Integration vs. Continuous Delivery vs. Continuous Deployment](#)



[Navigating the Continuous Delivery Pipeline](#)


[CONFERENCES](#)
[WEBINARS](#)
[VIRTUAL WORKSHOPS](#)
[ON-DEMAND](#)


## INTRODUCTION TO GITHUB ACTIONS - HOW TO EASILY AUTOMATE AND INTEGRATE WITH GITHUB

**VIRTUAL WORKSHOP**



BRENT LASTER  
@BRENTCLASTER

WORKSHOP STARTS

17 13 26 19  
DAYS HOURS MIN SEC

[REGISTER NOW](#)


[CONFERENCES](#)
[WEBINARS](#)
[VIRTUAL WORKSHOPS](#)
[ON-DEMAND](#)


## KUBERNETES FOR DEVELOPERS DEEP DIVE

**VIRTUAL WORKSHOP**

LIVE HANDS-ON FULL-DAY TRAINING  
STAY ENGAGED & STAY CURRENT

FRIDAY, AUGUST 5

11:00 AM EDT / 8:00 AM PDT



BRENT LASTER  
@BRENTCLASTER

WORKSHOP STARTS

24 13 25 10  
DAYS HOURS MIN SEC

[REGISTER NOW](#)



# Webinar sponsored by Uberconf

Getting Started

https://uberconf.com

HOME SCHEDULE SPEAKERS TRAVEL CONTACT US MEMBERS REGISTER

**Überconf**  
THE ULTIMATE CONFERENCE FOR SOFTWARE DEVELOPERS AND ARCHITECTS

July 18 - 21, 2023 — Denver, CO

[DAYS] 59 [HOURS] 21 [MINUTES] 05 [SECONDS] 59

Buy 4 Tickets, 5th is Free!

RECEIVE A \$500 AMAZON GIFT CARD

amazon

OFFER DETAILS

July 18 - 21, 2023

Denver, CO

REGISTER NOW

ÜberConf is our flagship educational event for software engineers and technical leaders. With 10 concurrent tracks, 140 sessions, and 10 full-day workshops; ÜberConf is truly the ultimate destination for passionate technologists.

Topics include: Software Architecture · Architecture · Modern Java · Kotlin · Leadership · Security · Front-End · DevOps · Microservices · Kubernetes · Docker · AI · Machine Learning · Gradle · Android · Performance · RxJava

# Containers Overview





# What are Containers?

- A container is a standard unit of software that functions like a fully provisioned machine installed with all the software needed to run an application.
- It's a way of packaging software so that applications and their dependencies have a self-contained environment to run in, are insulated from the host OS and other applications - and are easily ported to other environments.
- A container is NOT a VM. A container leverages several features of the Linux OS to "carve out" a self-contained space to run in.

What's in a container?



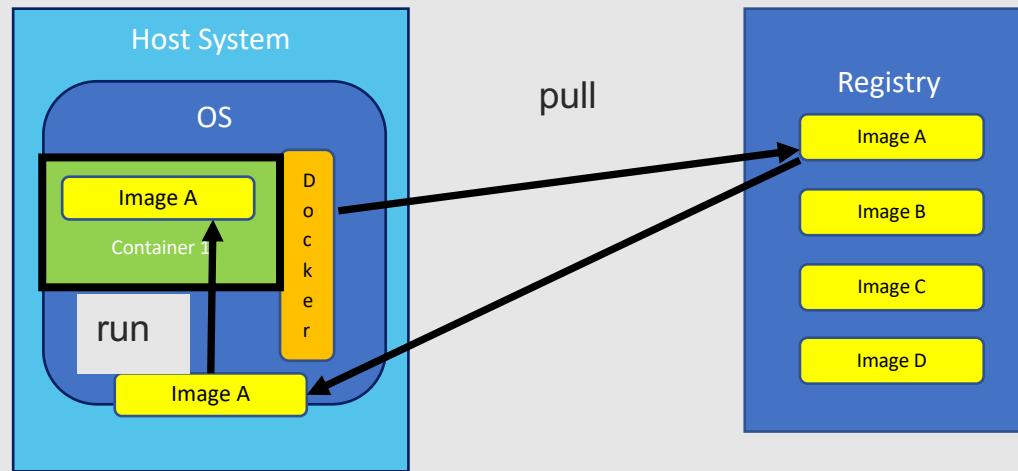
**Containers are running instances of images  
Images define what goes into a container.  
Containers are built from images.**

## Benefits of containers



# Easy to startup

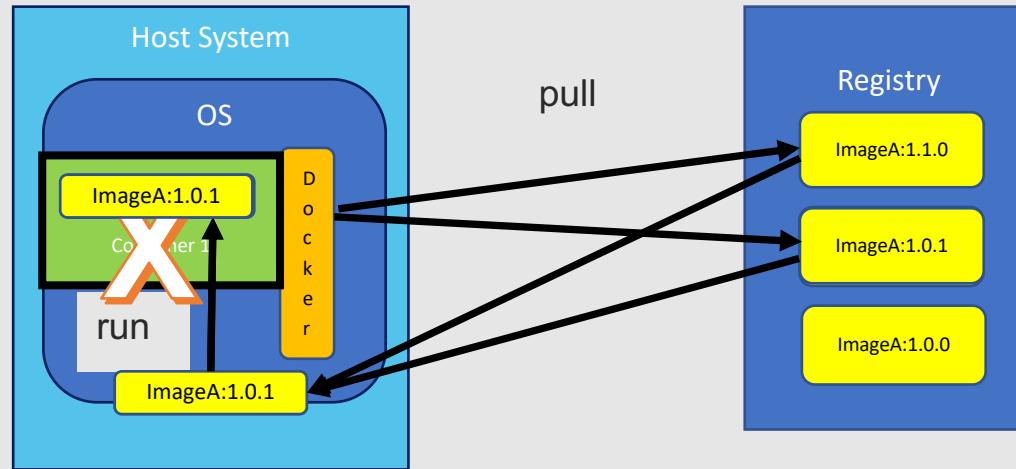
- Get image
- Run image as a container





# Quick, easy rollbacks

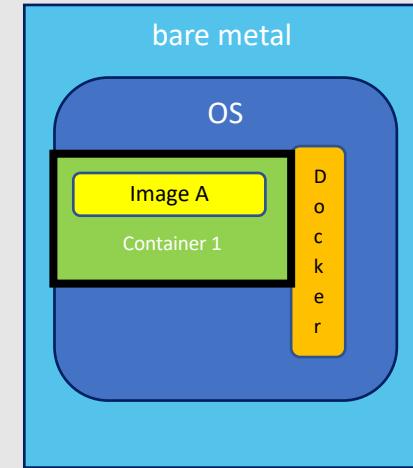
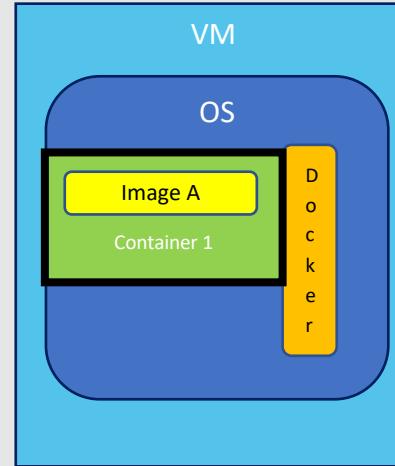
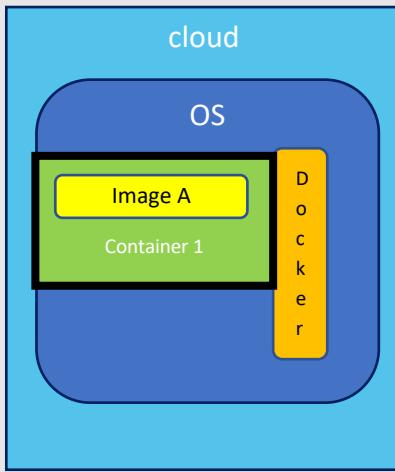
- Get image
- Run image as a container





# Portability

- Run nearly anywhere without rebuilding
- Works same way since container is the environment

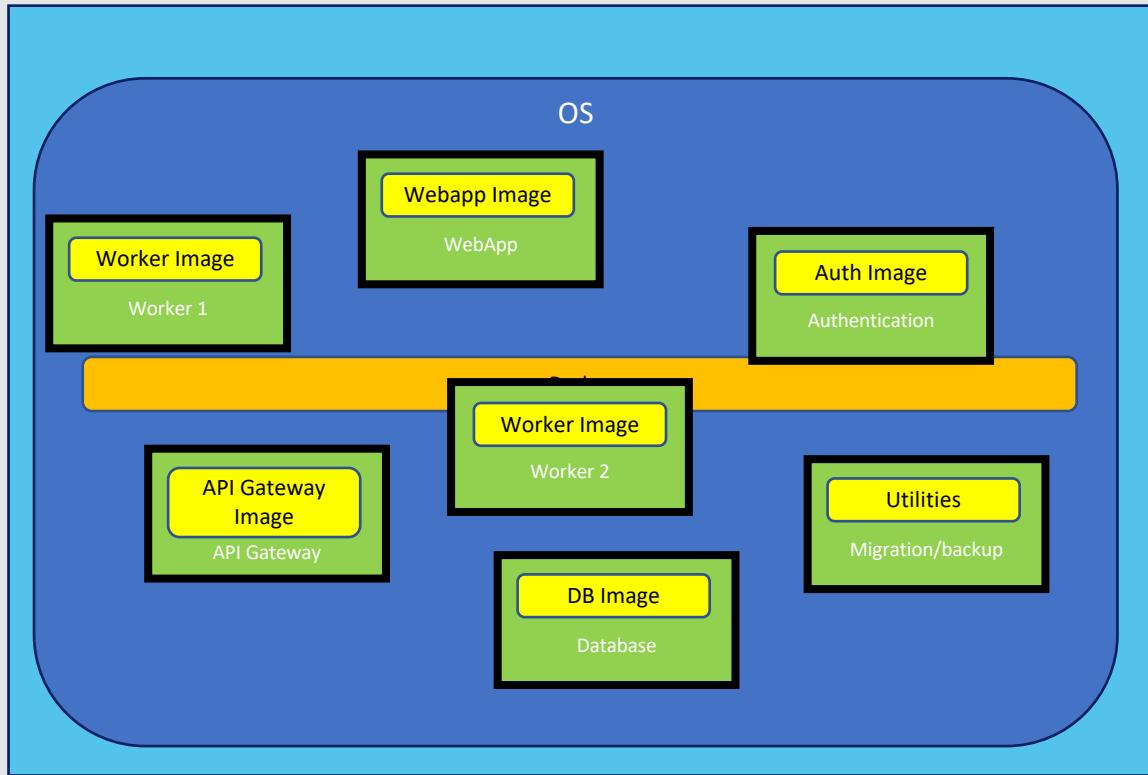




# Supports Distributed/Decoupled Architecture

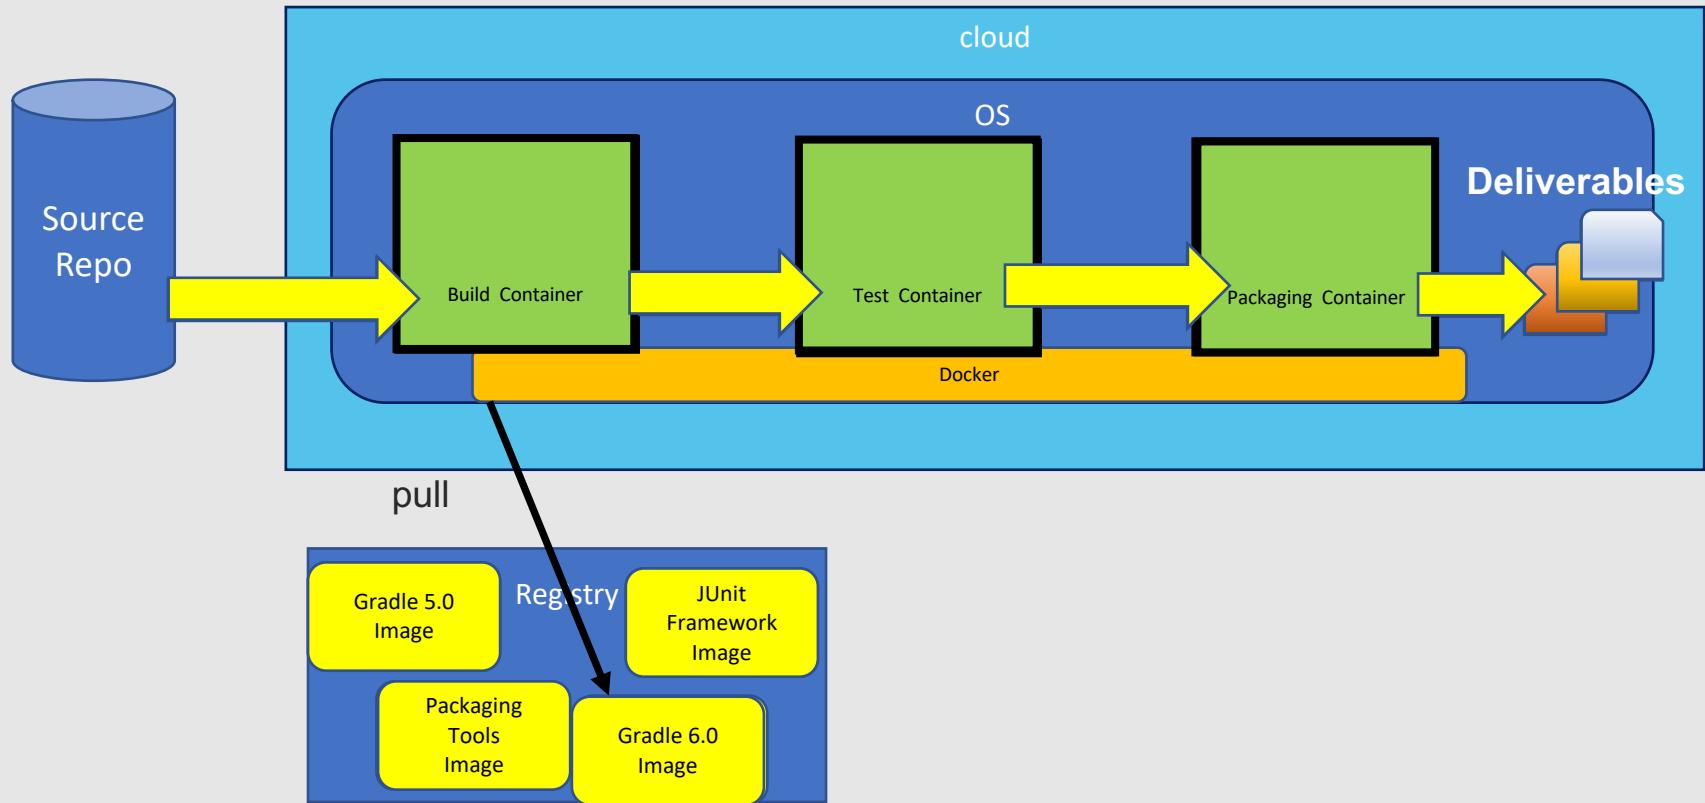
18

- **Loosely coupled services**
- **Removes need for monolithic stacks**





- Define containers well-suited for a specific task
- Spin up/down containers quickly/easily to address needs
  - such as devops workflow





# How do we work with multiple containers?

- Docker compose tool allows us to start containers together
- Uses service specifications in docker-compose.yml file
- Simplifies running multiple containers that need to be “linked” together
  
- Composing containers together is essentially a 3-step process
  - Create needed Dockerfiles for each part of app
  - Define services that make up the app in a docker-compose.yml file
  - Run docker-compose up and it starts and runs your app
  
- Docker run
  - Could also use docker run command in extended form with –link option
  
- Docker stack
  - Command in Docker CLI - lets you manage cluster of containers via Swarm (competitor to Kubernetes)
  - Can also leverage docker-compose.yml file
  
- Kubernetes or Kubernetes Environments

```

1  roar-web-container:
2    image: roar-web
3    ports:
4      - "8089:8080"
5    links:
6      - "roar-db-container:mysql"
7  roar-db-container:
8    image: roar-db
9    ports:
10   - "3308:3306"
11   environment:
12     MYSQL_USER: admin
13     MYSQL_PASSWORD: admin
14     MYSQL_DATABASE: registry
15     MYSQL_ROOT_PASSWORD: root+1

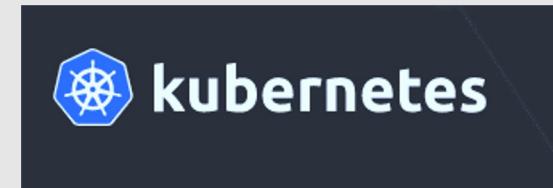
```

## Kubernetes Overview



# What is Kubernetes?

- A portable, extensible platform for managing containerized workloads and services (cluster orchestration system)
- Name derived from Greek for helmsman, thus the icon
- Frequently abbreviated as “k8s”
- Formerly known as “borg” – an internal Google project
- Open-sourced by Google in 2014
- Groups containers that make up an application into logical units for easy management and discovery.
- Goal is to provide a robust platform for running many containers.
- Allows automation of deployment, scaling, and managing containerized workloads.
- Kubernetes provides you with a framework to run distributed systems (of containers) resiliently.
- Takes care of
  - scaling requirements
  - failover
  - deployment patterns





# So how do we think about this?

- Analogy: Datacenter for containers
  - If we think of images/containers as being like computers we stage and use
  - We can think of Kubernetes as being like a datacenter for those containers
  - Main jobs of datacenter
    - » Provide systems to service needs (regardless of the applications)
    - » Keep systems up and running
    - » Add more systems / remove systems depending on load
    - » Deal with systems that are having problems
    - » Deploy new systems when needed
    - Provide simple access to pools of systems
    - Etc..





# Kubernetes is everywhere

- Has effectively “won the war” over other competitors
  - Docker swarm
  - Mesos
- Cloud providers all endorse it and provide ways to get a Kubernetes cluster
- Implementations on other enterprise platforms



# Kubernetes Features

25

- Service discovery and load-balancing
- Automated rollouts and rollbacks
- Storage orchestration
- Batch execution
- Self healing
- Secret and configuration management
- Horizontal scaling
- Automatic bin packing



# K8s Quick Terminology

- Cluster - an HA set of computers coordinated by k8s to work as a unit.
- Pods – object that contains and manages one or more containers and any attached volumes
- Service – abstraction that groups together pods based on identifiers called labels (or other characteristic)
- Deployment – defines a stateless app with a set number of pod replicas (scaled instances)
- Ingress – resource that lets cluster applications be exposed to external traffic
- Namespace - a logical area that groups k8s items like pods



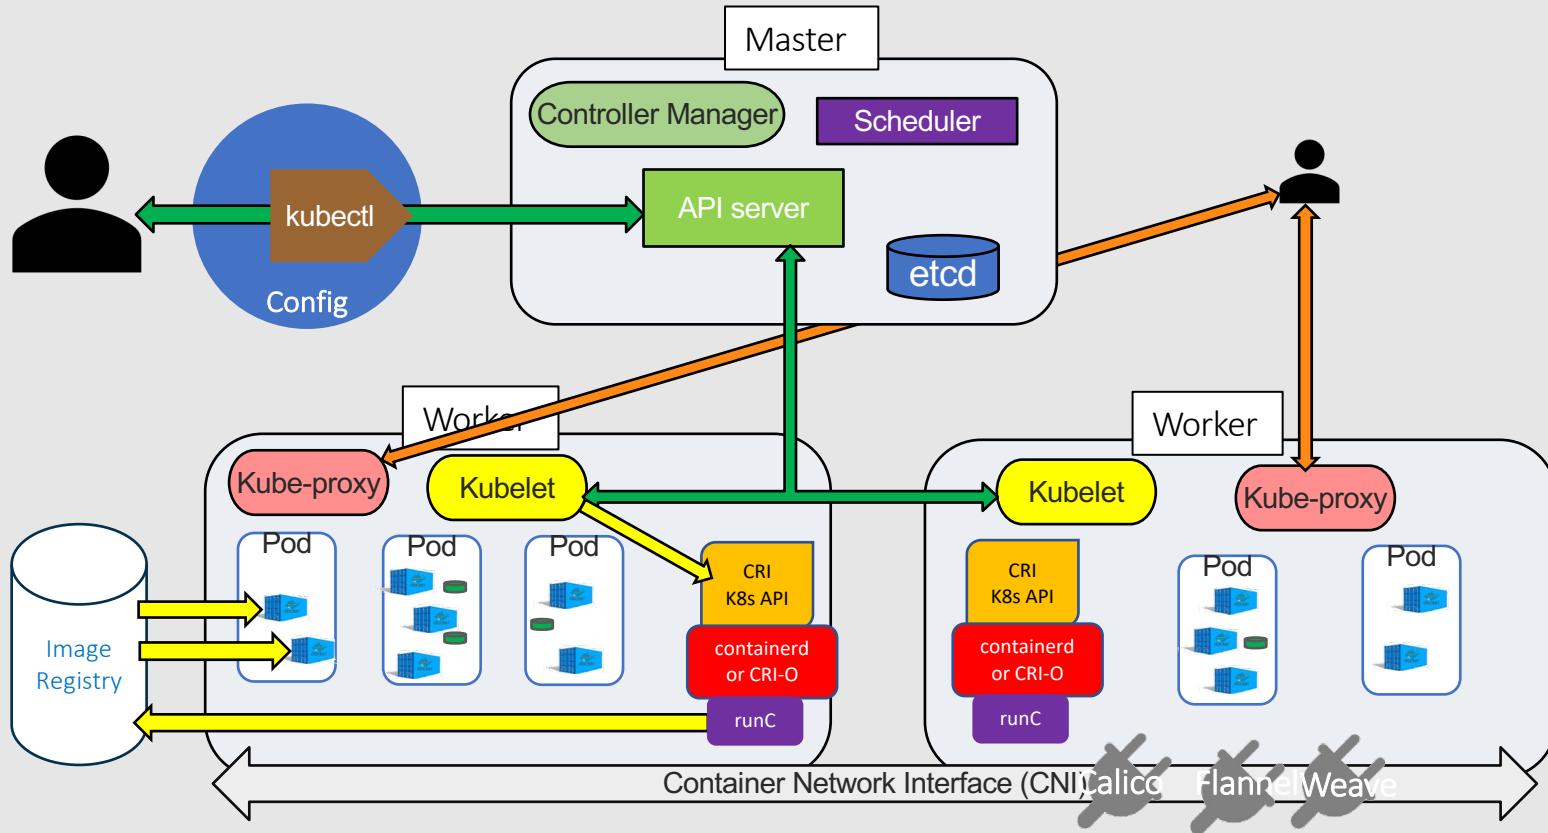
# Cluster Overview

cluster = HA set of computers coordinated by K8s to work as a unit

K8s abstractions allow deploying containers w/o requiring tying to specific host

K8s automates distribution & scheduling of containers efficiently

Nodes are either master (coordinate the work) or worker (run the applications)



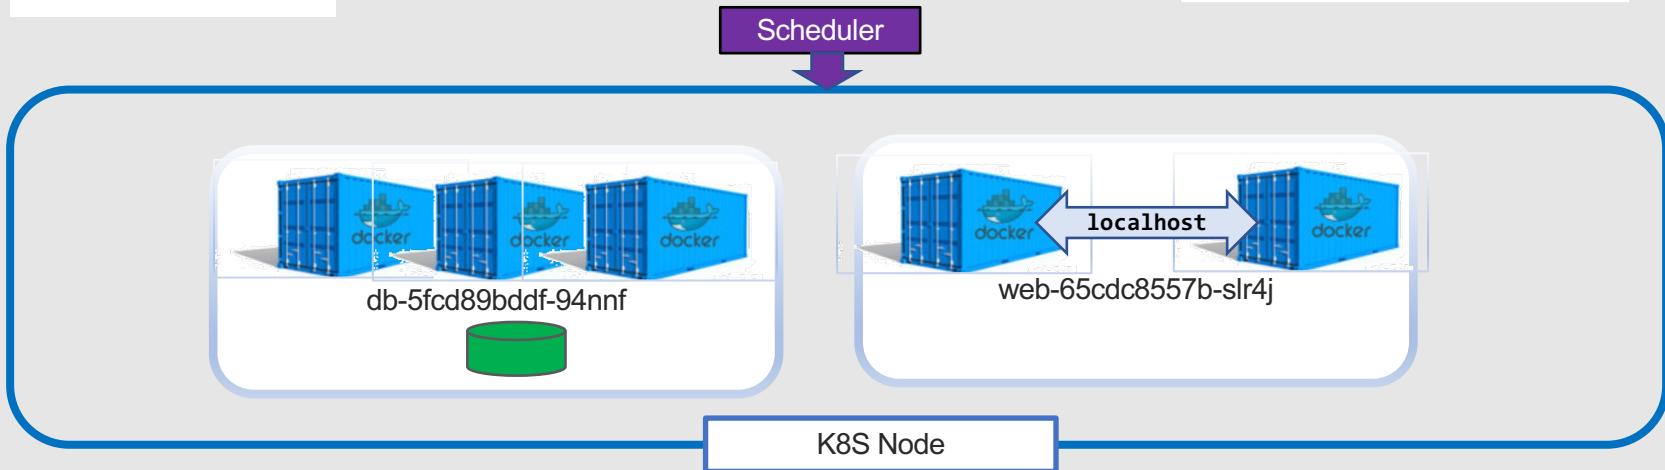


# Pods

smallest  
deployable  
unit in  
Kubernetes

represents a group of  
1 or more containers  
and any shared  
resources, such as  
storage volumes

scheduled on nodes  
by scheduler  
process



nodes in same pod can  
also talk to each other  
over "localhost"

containers should only  
be scheduled together  
if tightly coupled and  
need to share resources

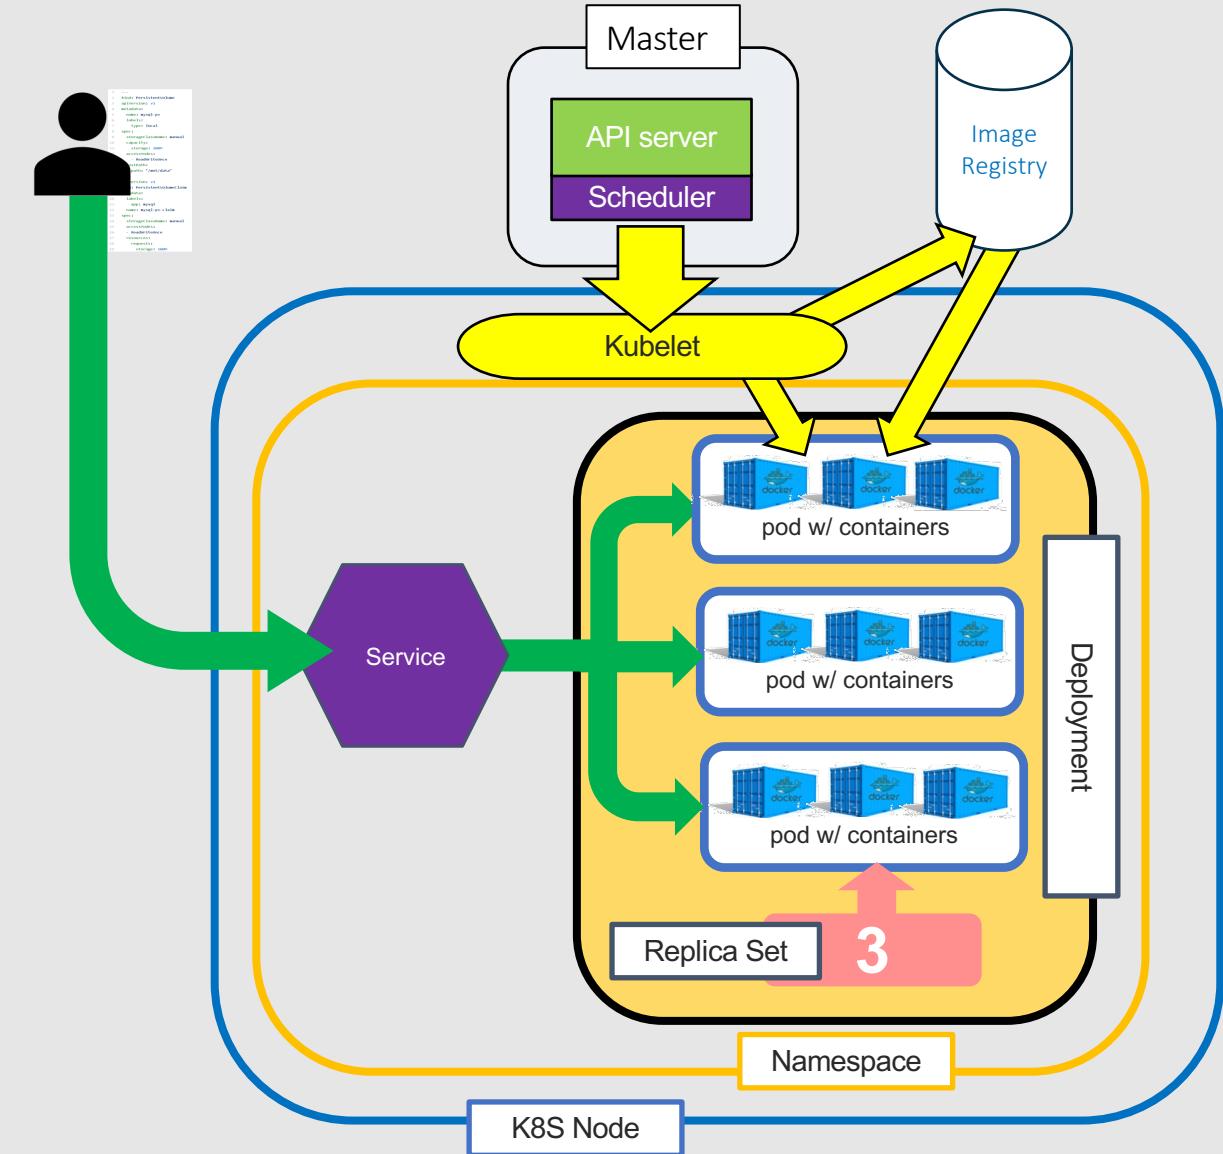
can use node  
"selector" to  
specify node to run  
on



# How are containers organized on K8s?

29

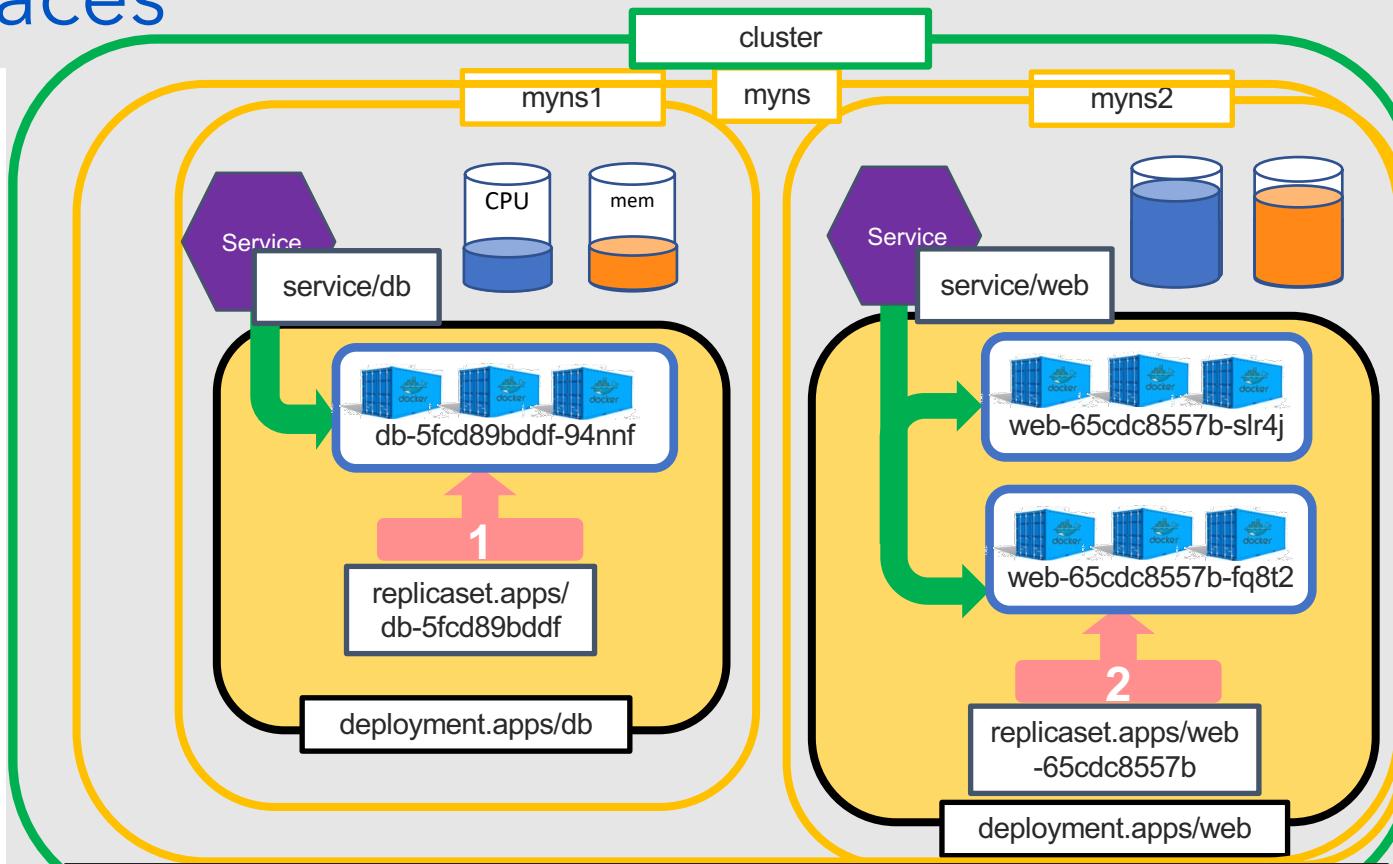
- K8s clusters have nodes
- Manifests (specs) define the objects
- Containers are created from images
- Pods wrap containers
- Pods are scaled (replicated) by deployments
- Pods are front-ended by services
- Namespaces group objects like pods, services, deployments together





# Namespaces

- Logical, unique working area within a cluster
- Partitions cluster's resources
- Resources are associated/scoped to a namespace
- Namespaces can have assigned quotas for resource use/limits
- All resources within a namespace must be unique in name (Kubernetes can generate unique suffixes for similar resources)
- Namespace "default" is default
- Default namespace is the one used if another is not specified
- Object is abbreviated as "-n" when filtering by one
- Can set context to change current one



```
$ kubectl get pods -n myns
NAME           READY   STATUS    RESTARTS   AGE
db-5fcdbbddf-94nnf   1/1    Running   0          4h18m
web-65cdc8557b-slr4j   1/1    Running   0          4h1m
web-65cdc8557b-fq8t2   1/1    Running   0          4h1m
```

NAME	READY	STATUS	RESTARTS	AGE
pod/db-5fcdbbddf-94nnf	1/1	Running	0	4h23m



# Kubectl Command Line

- Kubectl is the command line interface for running commands against k8s clusters
- Different pronunciations – usually either “kube control” or “kube cuttle” or “kube cuddle”
- Uses a “kube config” file to understand which cluster to work against and any necessary configuration information (kube config file may be sharable to other instantiations)
- Basic syntax is : kubectl [command] [TYPE] [NAME] flags
  - Examples:
    - **kubectl get** - get basic information about existence of objects
    - **kubectl apply** – create/update applications from files defining resources
    - **kubectl create** - create a new object
    - **kubectl describe** – get detailed info about current state of an object
    - **kubectl delete** – delete an object
- Resources can be referred to with resource/name syntax
  - kubectl describe pod/mypod-name
- Options can be in different locations
  - kubectl -n <namespace> get all
  - kubectl get all –n <namespace>
- Can use different forms of resource type names
  - singular, plural, abbreviations (ex: pod, pods, po)
- Abbreviations exist for most resources
  - po = pod, svc = service, deploy = deployment, etc.

```
diyuser3@training1:~$ kubectl get service -n roar2
NAME        TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)
mysql       ClusterIP   10.109.196.123  <none>        3306/TCP
roar-web    NodePort    10.106.132.164  <none>        8089:30318
diyuser3@training1:~$ kubectl -n roar2 get svc
NAME        TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)
mysql       ClusterIP   10.109.196.123  <none>        3306/TCP
roar-web    NodePort    10.106.132.164  <none>        8089:30318
diyuser3@training1:~$ kubectl get -n roar2 all | grep service
service/mysql     ClusterIP   10.109.196.123  <none>
service/roar-web  NodePort    10.106.132.164  <none>
```





# Configs and Contexts

32

- kubeconfig files
  - organizes info about clusters, users, namespaces, and authentication mechanisms
  - used by kubectl to find info needed to pick cluster and talk to cluster's API server
  - default name/location is \$HOME/.kube/config
  - can override default with KUBECONFIG env var or --kubeconfig flag on command line
- context
  - in kubeconfig file, groups together access parameters under name
  - access parameters are cluster, namespace (optional), and user
  - command line access to contexts is through "kubectl config <context>"

```
% kubectl config
Modify kubeconfig files using subcommands like "kubectl config set
current-context my-context"

Available Commands:
  current-context Display the current-context
  delete-cluster  Delete the specified cluster from the kubeconfig
  delete-context  Delete the specified context from the kubeconfig
  delete-user     Delete the specified user from the kubeconfig
  get-clusters    Display clusters defined in the kubeconfig
  get-contexts   Describe one or many contexts
  get-users      Display users defined in the kubeconfig
  rename-context Rename a context from the kubeconfig file
  set             Set an individual value in a kubeconfig file
  set-cluster    Set a cluster entry in kubeconfig
  set-context    Set a context entry in kubeconfig
  set-credentials Set a user entry in kubeconfig
  unset          Unset an individual value in a kubeconfig file
  use-context    Set the current-context in a kubeconfig file
  view           Display merged kubeconfig settings or a specified
  kubeconfig     file

Usage:
  kubectl config SUBCOMMAND [options]
```



# YAML and K8S specifications

- YAML is a type of markup language to define Kubernetes specs for resources
- Stored in .yaml or .yml text file
- Kubectl apply can take such a file as input and update cluster based on specs
  - Turns yaml specs into resources/objects running in cluster
- Kubectl get –o yaml can be used to dump out spec and status as yaml from running object
- Conventional block format uses a hyphen + space to denote a new item in a list
- Keys are separated from values by a colon + space; indented blocks use indentation and newlines to separate key-value pairs
- Strings do not (generally) require quotation marks
- Data structure hierarchy is maintained by outline indentation

```
--- # Favorite movies
- Casablanca
- North by Northwest
- The Man Who Wasn't There
```

```
--- # Indented Block
name: John Smith
age: 33
receipt: Oz-Ware Purchase Invoice
date: 2012-08-06
customer:
  first_name: Dorothy
  family_name: Gale

items:
  - part_no: A4786
    descrip: Water Bucket (Filled)
    price: 1.47
    quantity: 4

  - part_no: E1628
    descrip: High Heeled "Ruby" Slippers
    size: 8
    price: 133.7
    quantity: 1
```



# Understanding Kubernetes Objects

34

- To work with k8s objects, you use the k8s API
  - Kubectl command-line tool makes calls to API for you
  - Could also use k8s api client libraries
- K8s objects are persistent entities in the k8s system.
- K8s uses these entities to represent the state of the cluster.
  - They can describe:
    - What application containers are running and on which nodes.
    - Resources available to applications.
    - Policies around how those applications behave.
- Kubernetes object is “record of intent”
  - After creation, k8s will work to ensure object exists
  - Creating an object declares what you want cluster workload to look like
  - Known as cluster’s “desired state”
- Declarative model vs. imperative model



# Defining a Kubernetes Object in text

35

- When creating an object in k8s, have to provide
  - object spec to describe desired
  - basic info, such as a name
- K8s API expects info as JSON in body of request
- But, usually provide it from YAML file
- Kubectl command line converts to JSON for you
- Required fields
  - apiVersion – which version of the k8s API is being used to create the object
  - kind - what kind of object to create
  - metadata - data that helps uniquely identify the object, such as name, UID, namespace (optional)
  - Object's spec
    - Format different for every k8s object
    - Contains nested fields specific to the object
    - Can find details on specs in the API reference

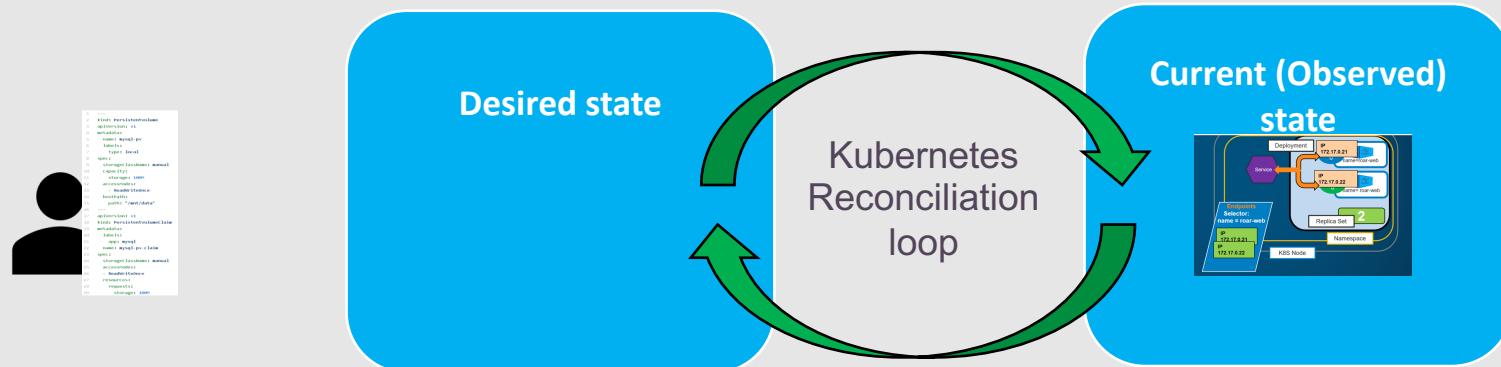
<https://kubernetes.io/docs/concepts/overview/working-with-objects/kubernetes-objects/>

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: roar-web
  labels:
    name: roar-web
  namespace: roar
spec:
  replicas: 1
  template:
    metadata:
      labels:
        name: roar-web
    spec:
      containers:
        - name: roar-web
          image: localhost:5000/roar-web:v1
          ports:
            - name: web
              containerPort: 8080
...
apiVersion: v1
kind: Service
metadata:
  name: roar-web
  labels:
    name: roar-web
  namespace: roar
spec:
  type: NodePort
  ports:
    - port: 8089
```



# Kubernetes is a Desired-State System

- User supplies desired state via declaring it in manifests
- Kubernetes works to balance the current state and the desired state
  - Desired state – what you want your production environment to be
  - Current (observed) state – current status of your production environment





# Kubernetes Objects - Spec and Status

- Every running k8s object includes two nested object fields
  - spec
    - User-provided
    - Describes desired state for the object
  - status
    - Provided and updated by k8s system
    - Describes the actual state of the object

```
diyuser3@training1:~$ kubectl edit -n roar2 pod roar-web-74bb47bdb8-56l4n
```

```

apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2019-07-21T20:04:34Z"
  generateName: roar-web-74bb47bdb8-
  labels:
    app: roar-web
    pod-template-hash: 74bb47bdb8
  name: roar-web-74bb47bdb8-56l4n
  namespace: roar2
  ownerReferences:
    - apiVersion: apps/v1
      blockOwnerDeletion: true
      controller: true
      kind: ReplicaSet
      name: roar-web-74bb47bdb8
      uid: 89d00855-abf0-11e9-b30c-080027c4188a
  resourceVersion: "346085"
  selfLink: /api/v1/namespaces/roar2/pods/roar-web-74bb47bdb8-56l4n
  uid: c2b50783-abf2-11e9-b30c-080027c4188a
spec:
  containers:
    - image: localhost:5000/roar-web:v2
      imagePullPolicy: Always
      name: roar-web
      ports:
        - containerPort: 8080
          name: web
          protocol: TCP
status:
  secret:
    defaultMode: 420
    secretName: default-token-dxx7t
  status:
    conditions:
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: Initialized
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: Ready
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: ContainersReady
      - lastProbeTime: null
        lastTransitionTime: "2019-07-21T20:04:34Z"
        status: "True"
        type: PodScheduled
    containerStatuses:
      - containerID: docker://41c2d31ef46b8ea1
        image: localhost:5000/roar-web:v2
        imageID: docker-pullable://localhost:5000/roar-web:v2
        lastState:
          terminated:
            containerID: docker://c0ddd9686971
            exitCode: 143

```



# kubectl edit

- Allows editing of resources in place with defined editor

Examples:

```
# Edit the service named 'docker-registry':
kubectl edit svc/docker-registry
```

```
# Use an alternative editor
KUBE_EDITOR="nano" kubectl edit
svc/docker-registry
```

```
# Edit the job 'myjob' in JSON using the v1
API format:
```

```
kubectl edit job.v1.batch/myjob -o json
```

```
# Edit the deployment 'mydeployment' in
YAML and save the modified config in
its annotation:
```

```
kubectl edit deployment/mydeployment -o
yaml --save-config
```

```

# Please edit the object below. Lines beginning with
# a '#' will be ignored,
# and an empty file will abort the edit. If an error
occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: apps/v1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
    meta.helm.sh/release-name: ts
    meta.helm.sh/release-namespace: ts
  creationTimestamp: "2021-06-22T22:08:53Z"
  generation: 2
  labels:
    app: mysql
    app.kubernetes.io/managed-by: Helm
    chart: roar-db-0.1.0
    release: ts
  name: mysql
  namespace: ts
  resourceVersion: "84294"
  uid: 5c0acb7d-8007-424a-8539-469be75f3166
spec:
  progressDeadlineSeconds: 600
  replicas: 1
  revisionHistoryLimit: 10
  selector:

```

```
$ k edit deploy/mysql
```

```
deployment.apps/mysql edited
```



# Labels

29

- Labels are the main mechanisms used in Kubernetes to select/organize/associate objects
- Label is a key-value pair without any pre-defined meaning
- Kubectl get has option “--show-labels” to show labels from objects

```
$ kubectl get pods -n istio1 --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
roar-current-6c75d49d78-bztfd	3/3	Running	6	6d21h	app=roar,version=current
roar-new-559997dc54-7qhjh	3/3	Running	6	6d21h	app=roar,version=new

```
apiVersion: v1
kind: Pod
metadata:
  name: roar-current
  labels:
    app: roar
    version: current
spec:
  containers:
    - image: localhost:5000/roar-web:v1
      imagePullPolicy: Always
      name: roar-web
      ports:
        - containerPort: 8080
          name: roar-web
          protocol: TCP
      resources: {}
```

- Can add label to object with “kubectl label” command

```
$ kubectl label deploy roar-new -n istio1 type=experimental
deployment.extensions/roar-new labeled
```

```
apiVersion: v1
kind: Pod
metadata:
  name: roar-new
  labels:
    app: roar
    version: new
spec:
  containers:
    - image: localhost:5000/roar-web:v2
      imagePullPolicy: Always
      name: roar-web
      ports:
        - containerPort: 8080
          name: roar-web
          protocol: TCP
      resources: {}
```

- We can filter based on a label
  - --selector is long form of option
  - -l is short form
- Most k8s objects support set-based selectors (choosing which items based on a set to select)
  - do:

```
$ kubectl get pods -n istio1 -l 'version in (current, new)'
```

NAME	READY	STATUS	RESTARTS	AGE
roar-current-6c75d49d78-bztfd	3/3	Running	6	6d21h
roar-new-559997dc54-7qhjh	3/3	Running	6	6d21h

- Labels can apply to other objects, such as nodes and services and be used in other operations

```
$ kubectl delete pods -l 'version in (current, new)'
```

- Types: “Equality-based” (=, !=) “Set-based” (in ())



# Deployments

- Spinning Up Pods

- A deployment is like a pod “supervisor”.
- After initiating a deployment, we have the deployment, the replicaset, and the pods
- Naming of the pods and replicaset are derived from the deployment name

```

apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: roar-web
  labels:
    name: roar-web
  namespace: roar
spec:
  replicas: 1
  template:
    metadata:
      labels:
        name: roar-web
    spec:
      containers:
        - name: roar-web
          image: localhost:5000/roar-web:v1
          ports:
            - name: web
              containerPort: 8080

```

```

diyuser3@training1:~$ k get deploy -n roar | grep web
roar-web   1/1     1       1           6d23h
diyuser3@training1:~$ k get rs  -n roar | grep web
roar-web-f95cf6574  1       1       1           6d23h
diyuser3@training1:~$ k get pod  -n roar | grep web
roar-web-f95cf6574-ndt4k  1/1     Running   0           9m44s

```



# Deployments

- Updating Pods

- What happens if we make a significant change (update image to v2) and apply it to the running instance?
- `kubectl apply -f roar-complete.yaml -n roar`
- Pods with old versions are terminated and new ones created with the updated version

```
$ kubectl get pods -n roar -w
```

NAME	READY	STATUS	RESTARTS	AGE
mysql-8559ccc47d-crvhv	1/1	Running	0	4h22m
roar-web-f95cf6574-wcm7z	1/1	Running	0	15s
roar-web-7bcc45c79-pp9hs	0/1	Pending	0	0s
roar-web-7bcc45c79-pp9hs	0/1	Pending	0	0s
roar-web-f95cf6574-wcm7z	1/1	Terminating	0	50s
roar-web-7bcc45c79-pp9hs	0/1	ContainerCreating	0	0s
roar-web-f95cf6574-wcm7z	0/1	Terminating	0	53s
roar-web-7bcc45c79-pp9hs	1/1	Running	0	3s

- And new replicaset is created

```
$ k get rs -n roar
```

NAME	DESIRED	CURRENT	READY	AGE
mysql-8559ccc47d	1	1	1	6d22h
roar-web-7bcc45c79	1	1	1	55s
roar-web-f95cf6574	0	0	0	7d5h

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: roar-web
  labels:
    name: roar-web
    type: experimental
  namespace: roar
spec:
  replicas: 1
  template:
    metadata:
      labels:
        name: roar-web
    spec:
      containers:
        - name: roar-web
          image: localhost:5000/roar-web:v2
          ports:
            - name: web
              containerPort: 8080
```

## Lab 1 - Exploring and Deploying into Kubernetes

**Purpose:** In this lab, we'll start to learn about Kubernetes and its object types, such as nodes and namespaces. We'll also deploy a version of our app that has had Kubernetes yaml files created for it. And we'll see how to do some simple debugging when Kubernetes deployments don't go as planned.

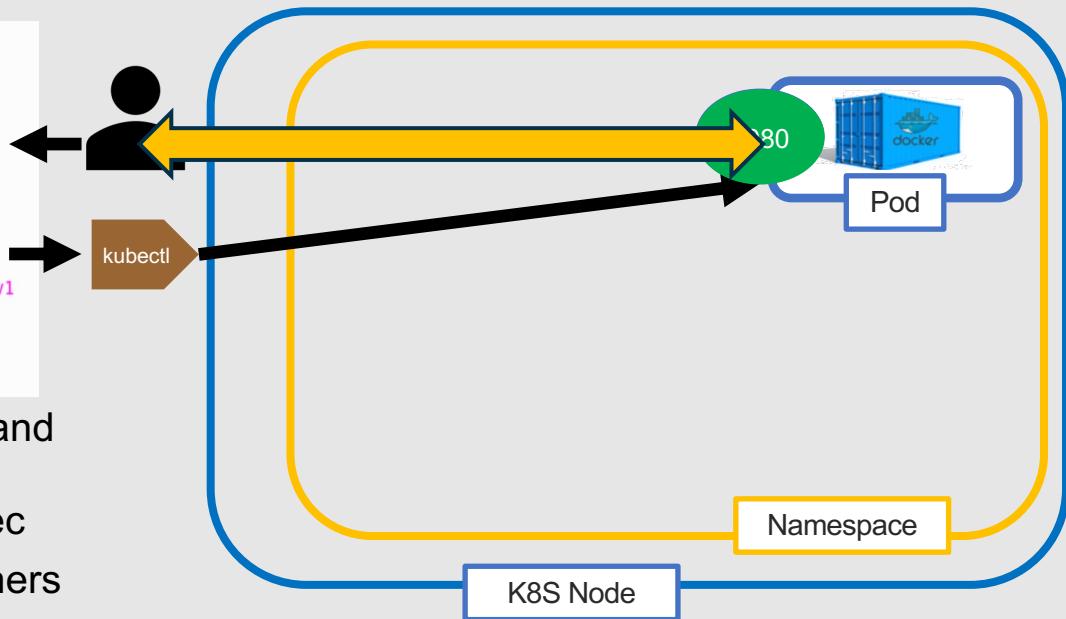


# Deployments vs. pods-only

- Using Pods without Deployments

```

apiVersion: v1
kind: Pod
metadata:
  name: roar-web
  labels:
    name: roar-web
  namespace: roar
spec:
  containers:
    - name: roar-web
      image: localhost:5000/roar-web:v1
      ports:
        - name: web
          containerPort: 8080
  
```

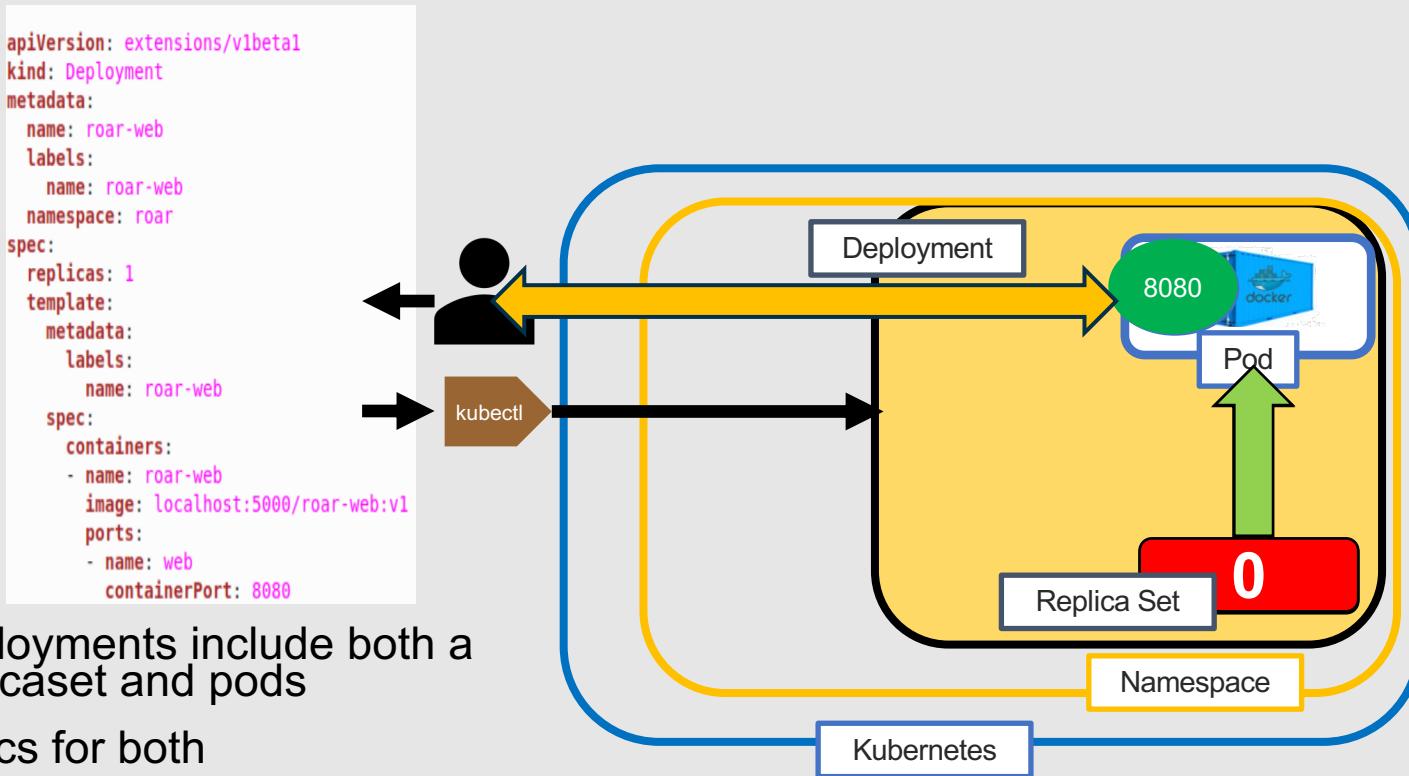


- Objects run within node and namespace
- Objects created from spec
- Pods encompass containers
- Ways to connect directly to them
- If pod goes away (and still need functionality), must manually start a new one
- Different ip address on new one



# Deployments vs. pods-only

- Pods with Deployments

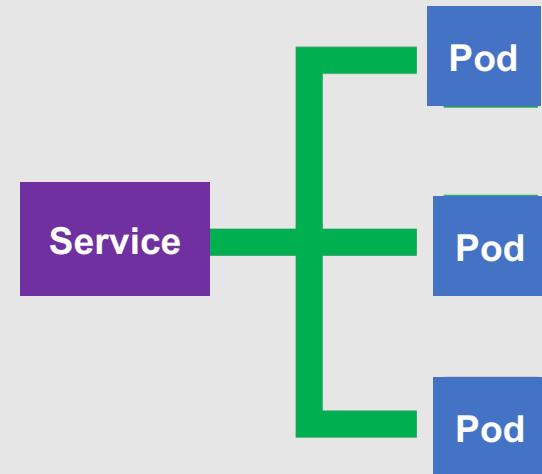


- Deployments include both a replicaset and pods
- Specs for both
- If pod goes away, replicaset starts up a new one automatically – up to number of replicas specified
- Different ip address on new one



# Services

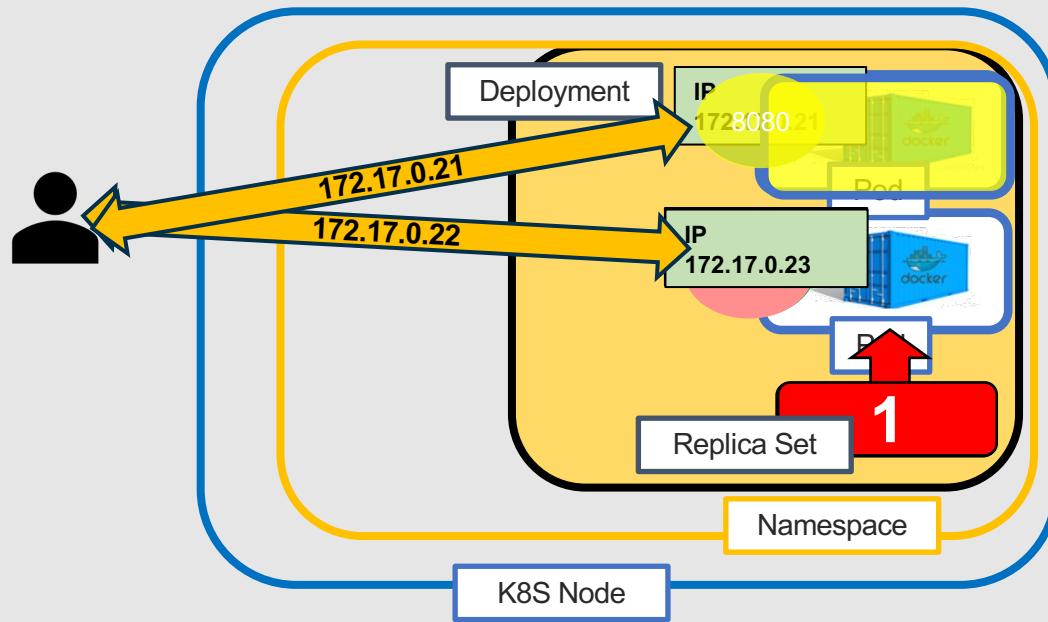
- A single virtual, stable connection for multiple volatile pods
  - Kubernetes service provides a level of abstraction for pods
  - Pods may be “volatile” so can’t rely on their respective IP addresses to connect to long-term
  - A service provides a virtual, consistent IP to connect to for a set of pods
  - Pods can be selected by labels
  - Virtual IP of a service is not connected to an actual network interface
  - Virtual IP exists to forward traffic to one or more pods
  - Kube-proxy process (running on every node) is responsible for keeping the mapping between the virtual IP and the pods up-to-date
  - Kube-proxy queries the API server to learn about new services





# Services

- Without Services
- We can connect to individual pods (if allowed)
- If in deployment and pod goes away, we lose access to that pod
- ReplicaSet in deployment will spin up another one but won't have same address
- Similar situation if a pod becomes unusable
  - may have other pods that could handle needs, but no way to automatically discover or connect

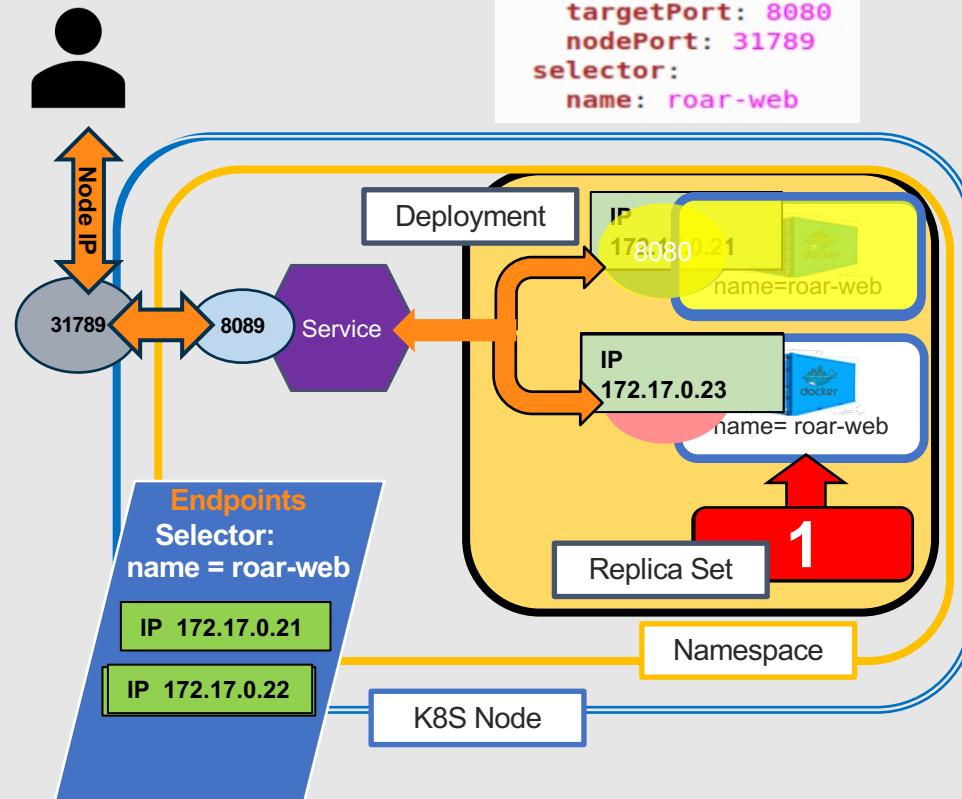




# Services

- With a service
- Service provides virtual address for us to connect to for frontend
- Uses labels to select “pool” of backend pods to map to
- Endpoints for a service are list of available ip addresses for usable pods
- Example service here is NodePort
- If one pod goes down or becomes unavailable, service can connect to another pod
- Spec is shown below

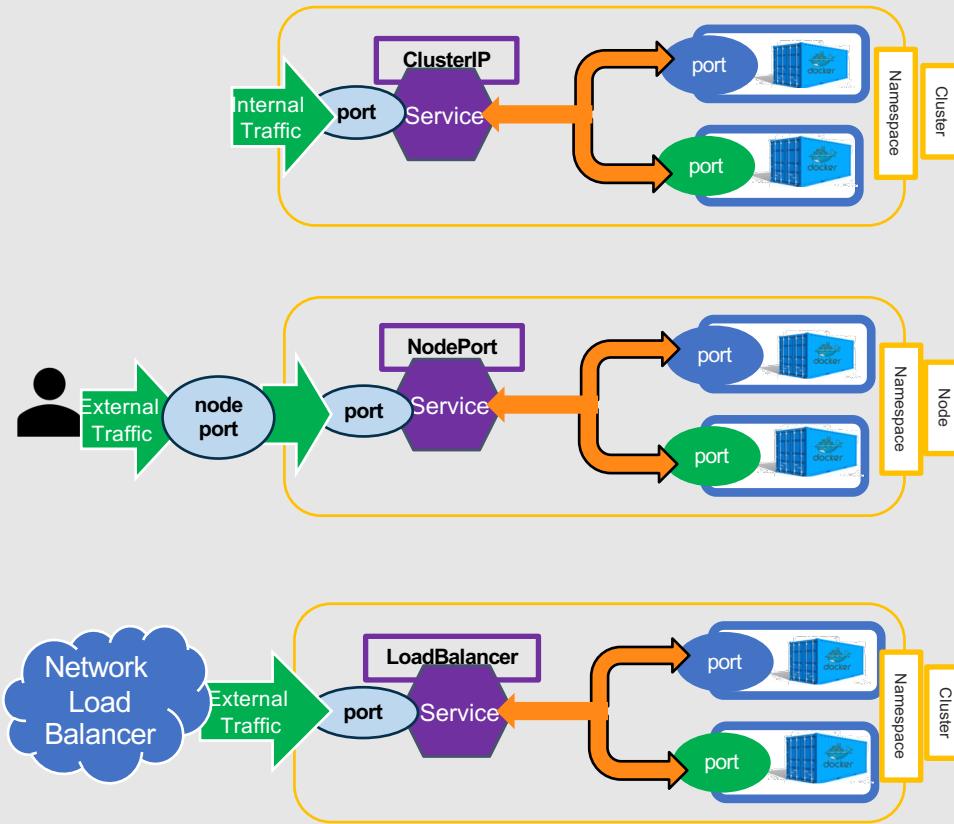
```
apiVersion: v1
kind: Service
metadata:
  name: roar-web
  labels:
    name: roar-web
  namespace: roar2
spec:
  type: NodePort
  ports:
    - port: 8089
      targetPort: 8080
      nodePort: 31789
  selector:
    name: roar-web
```





# Types of Services

- ClusterIP
  - Default K8S service
  - Provides service inside cluster for other apps in cluster to access
  - No external access (except via kubectl proxy for debugging, etc.)
- NodePort
  - Most basic way to get external traffic to service
  - Opens up a port on the K8S node
  - Any traffic going to that port on the node is forwarded to the service
  - Uses node IP address
  - Only ports in range 30000-32767
- LoadBalancer
  - Gives you Load Balancer with single IP that forwards all traffic to service
  - Default method for directly exposing a service
  - No filtering, no routing
  - Cost factor on clouds
- ExternalName
  - Maps service to contents of an external name field (foo.bar.com)
  - Returns a CNAME record with that value

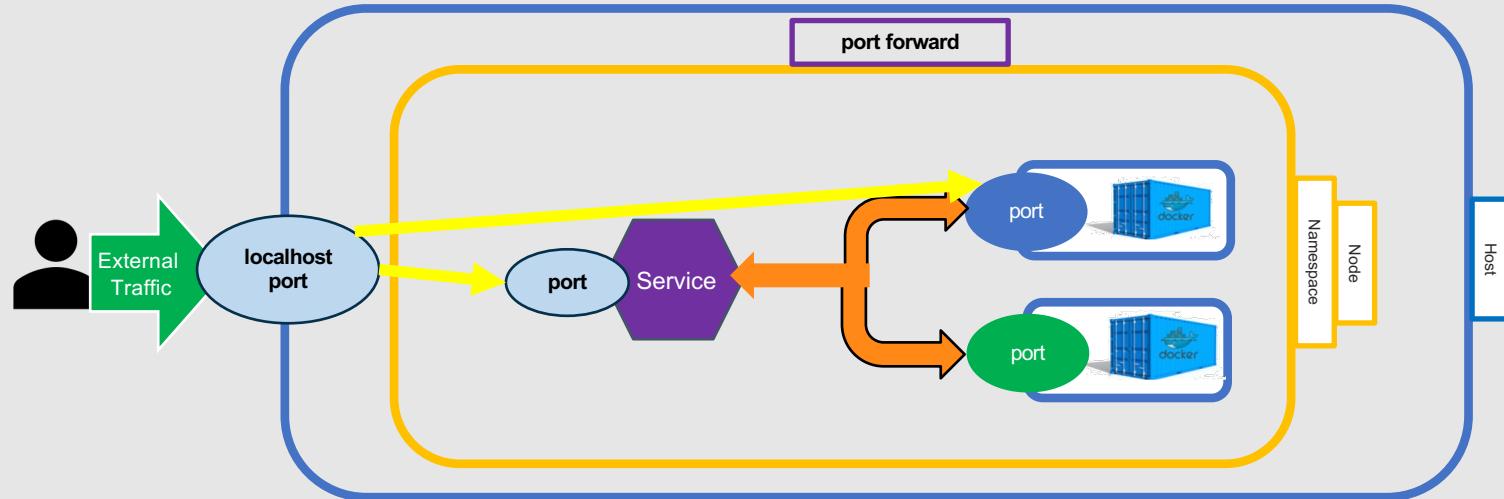




# Forwarding Ports in Kubernetes

49

- Allows you to access and interact with internal K8s cluster processes from your localhost
- Used to investigate and validate services locally
- Command in K8s is "kubectl port-forward"
  - Arguments include cluster resource name
  - port number to port-forward to
- Result is K8s API server establishes a single http connection between localhost and resource on the cluster
- Allows connection to any pod in cluster
  - if pod fails, connect and fix
- Especially useful for back-end services not usually intended for remote exposure



# Port Forward Examples

**Specific pod** - Selects specific pod with port number for container = 8080 and listening port = 30001

```
$ kubectl port-forward pod/roar-web-f5cbff465-vvr8f 30001:8080
```

**Random local port** - Listen on a random port locally, and forward to port 8089 within pod

```
$ kubectl port-forward svc/roar-web :8089
```

**Corresponding local and remote port** - Listen and forward data using identical port (8080) on localhost and in pod

```
$ kubectl port-forward pod/roar-web-f5cbff465-v8f 8080
```

**Specific Local IP Address for Port Forwarding** - Listen on port 8080 on localhost using defined IP; forward to port 5762 in the pod

```
$ kubectl port-forward --address 10.0.2.15 pod/roar-web-f5cbff465-vvr8f 8080
```

**Use deployment to select port-forward pod** - Listen and forward data using the same port 8080 both locally and within the pod. The deployment defines which pod to use

```
$ kubectl port-forward deploy/roar-web 8080
```

**Allow service to define port-forward pod** - Listen and forward data using the same ports (8080 5762) both locally and within the pod. The service selects which pod to use

```
$ kubectl port-forward svc/roar-web 8089
```



# Ingresses

51

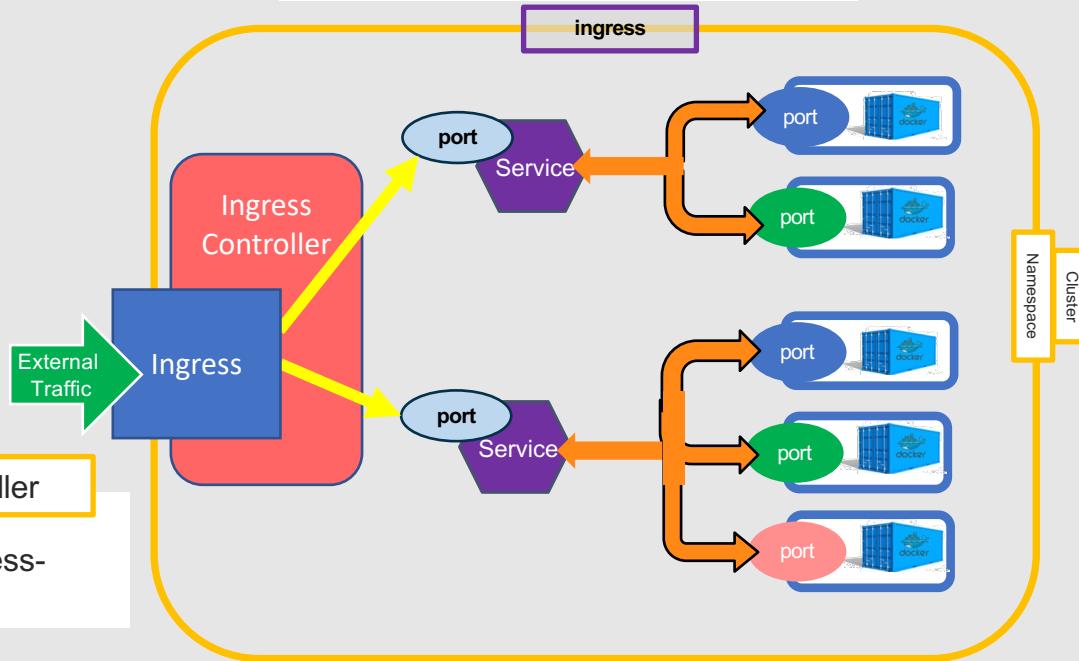
Ingress  
YAML

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
    paths:
      - path: /current
        backend:
          serviceName: v1-service
          servicePort: 8089
      - path: /beta
        backend:
          serviceName: v2-service
          servicePort: 8089
```

- Allows you to expose services outside of cluster
- Allows consolidation/configuration of routing rules in a single resource
- NodePort and LoadBalancer let you expose a service via the manifest
- Ingress is completely independent resource
- Declared, created, and destroyed independently
- Must have ingress controller for cluster

Install Ingress Controller

```
kubectl apply -f
https://raw.githubusercontent.com/kubernetes/ingress-nginx/master/deploy/provider/cloud-generic.yaml
```





# kubectl logs

Print the logs for a container in a pod or specified resource.

Examples:

```
# Return snapshot logs from pod nginx with only one container
kubectl logs nginx
```

```
# Return snapshot logs from all containers in pods defined by
label app=nginx
kubectl logs -lapp=nginx --all-containers=true
```

```
# Return snapshot of previous terminated ruby container
logs from pod web-1
kubectl logs -p -c ruby web-1
```

```
# Begin streaming the logs of the ruby container in pod web-1
kubectl logs -f -c ruby web-1
```

```
# Display only the most recent 20 lines of output in pod nginx
kubectl logs --tail=20 nginx
```

```
# Show all logs from pod nginx written in the last hour
kubectl logs --since=1h nginx
```

```
# Show logs from a kubelet with an expired serving certificate
kubectl logs --insecure-skip-tls-verify-backend nginx
```

```
# Return snapshot logs from first container of a job named hello
kubectl logs job/hello
```

```
# Return snapshot logs from container nginx-1 of a deployment
named nginx
kubectl logs deployment/nginx -c nginx-1
```

```
$ k logs mysql-5c97d489dd-z5pp2
```

```
2021-06-22 23:43:46+00:00 [Note] [Entrypoint]:
Entrypoint script for MySQL Server 5.7.32-
1debian10 started.
```

```
2021-06-22 23:43:46+00:00 [Note] [Entrypoint]:
Initializing database files
```

```
2021-06-22T23:43:46.950512Z 0 [Warning] TIMESTAMP
with implicit DEFAULT value is deprecated. Please
use --explicit_defaults_for_timestamp server
option (see documentation for more details).
```

```
2021-06-22T23:43:47.138897Z 0 [Warning] InnoDB:
New log files created, LSN=45790
```

```
2021-06-22T23:43:47.211492Z 0 [Warning] InnoDB:
Creating foreign key constraint system tables.
```

```
2021-06-22T23:43:47.270535Z 0 [Warning] No
existing UUID has been found, so we assume that
this is the first time that this server has been
started. Generating a new UUID: b02442ec-d3b3-
11eb-9519-0242ac110006.
```

```
2021-06-22T23:43:47.274194Z 0 [Warning] Gtid
table is not ready to be used. Table
'mysql.gtid_executed' cannot be opened.
```

```
2021-06-22T23:43:47.625895Z 0 [Warning] CA
certificate ca.pem is self signed.
```

```
2021-06-22T23:43:47.848450Z 1 [Warning]
root@localhost is created with an empty password
! Please consider switching off the --initialize-
insecure option.
```

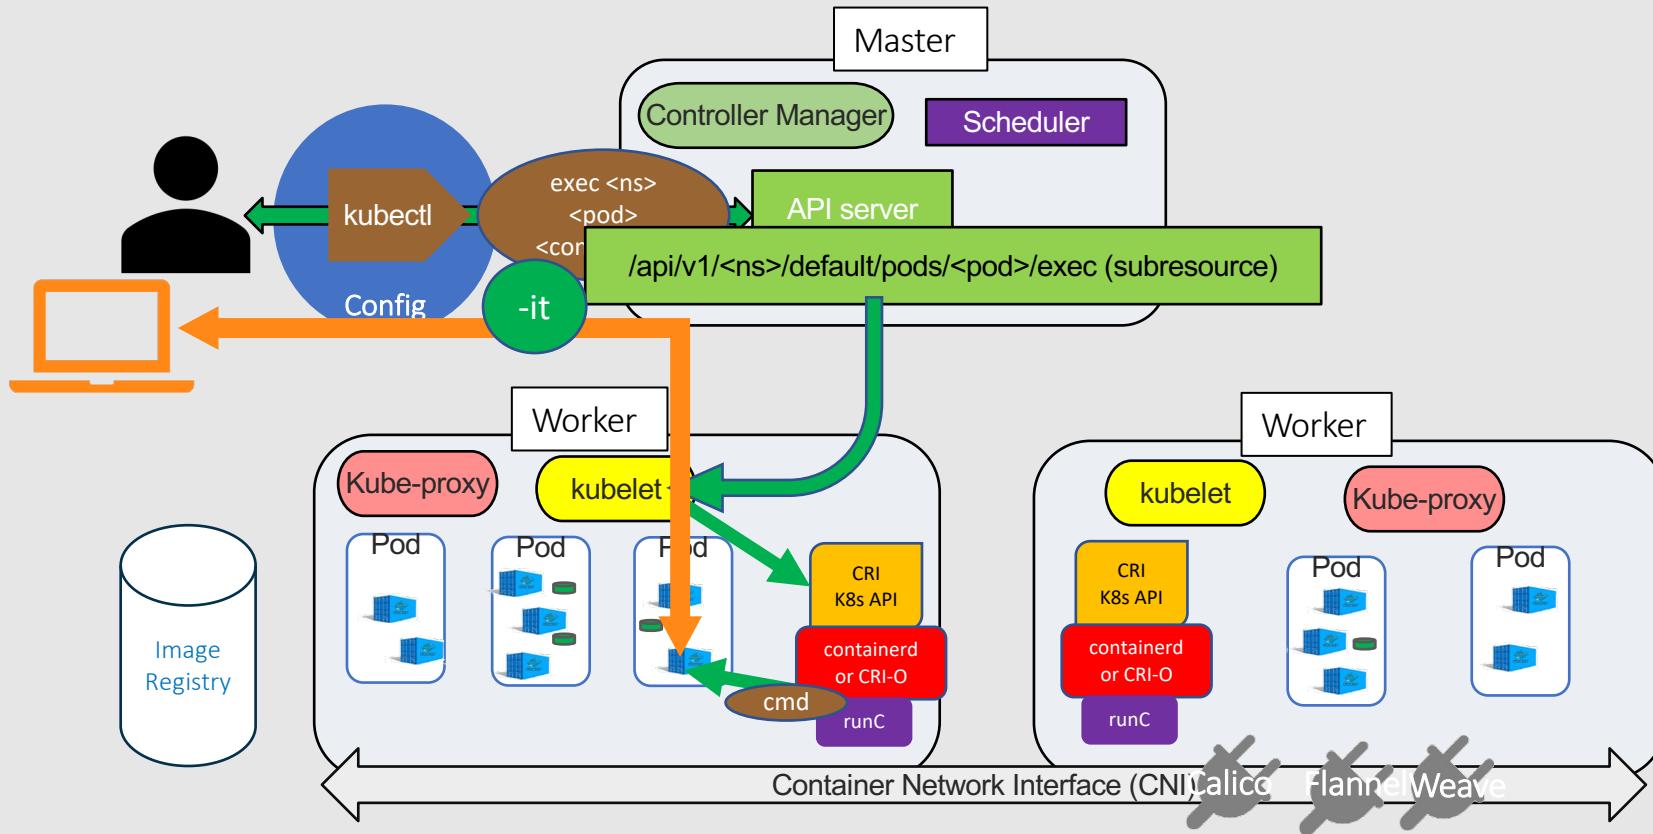
```
2021-06-22 23:43:52+00:00 [Note] [Entrypoint]:
Database files initialized
```

```
2021-06-22 23:43:52+00:00 [Note] [Entrypoint]:
Starting temporary server
```



# kubectl exec

- Allows for inspecting and debugging applications
- Allows executing commands inside our containers
- Can be interactive



## Lab 2 - Working with services and ports

Purpose: In this lab, we'll learn more about how Kubernetes services and ports work and how to debug basic issues with them.



# Data Center Analogy

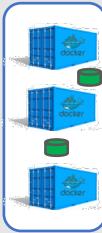
55

- Functions: Uptime, scaling, redundancy...

- Container in a pod ~ server in a rack



- Pod ~ rack of servers



- Deployment ~ multiple racks (replicas)



- Service ~ central control / login server



- Namespace ~ server room



55



# kubectl describe

# Describe a node

```
kubectl describe nodes kubernetes-node-emt8.c.myproject.internal
```

# Describe a pod

```
kubectl describe pods/nginx
```

# Describe a pod identified by type and name in "pod.json"

```
kubectl describe -f pod.json
```

# Describe all pods

```
kubectl describe pods
```

# Describe pods by label name=myLabel

```
kubectl describe po -l name=myLabel
```

# Describe all pods managed by the 'frontend' replication controller  
(rc-created pods)

# get the name of the rc as a prefix in the pod the name).

```
kubectl describe pods frontend
```

```
$ kubectl describe pod mysql-7bf6b7fc5-2tdkw
Name:      mysql-7bf6b7fc5-2tdkw
Namespace:  ts
Priority:   0
Node:       <none>
Labels:     app=mysql
            pod-template-hash=7bf6b7fc5
Annotations: <none>
Status:    Pending
IP:
IPs:      <none>
Controlled By: ReplicaSet/mysql-7bf6b7fc5
Containers:
  roar-db:
    Image:   quay.io/techupskills/roar-db:1.0.1
    Port:    3306/TCP
    Host Port: 0/TCP
    Limits:
      cpu:  1
      memory: 10Gi
    Requests:
      cpu:  1
      memory: 10Gi
    Readiness: exec [mysql] delay=5s timeout=1s period=10s #success=1 #failure=3
    Environment:
      MYSQL_DATABASE: registry
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-j2lcx (ro)
Conditions:
  Type        Status
  PodScheduled  False
Volumes:
  kube-api-access-j2lcx:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:  kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:   true
    QoS Class:    Guaranteed
    Node-Selectors: <none>
    Tolerations:   node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                  node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type  Reason     Age           From           Message
  ----  ----     --           --           --
  Warning FailedScheduling  15s (x23 over 42m)  default-scheduler  0/1 nodes are available:
  1 Insufficient memory.
```



# Configmaps and Secrets

57

- Approaches for configuring applications
- Could specify environment variables in Dockerfile or k8s yaml file
- Bad part of environment variables is they are tied to container or deployment.
- To change them, must modify container or redo deployment.
- Data must also be duplicated across objects.
- Alternative is secrets for confidential data and configmaps for non-confidential data
- Difference between Secrets and ConfigMaps are that Secrets are obfuscated with a Base64 encoding.
- Can be created like any other objects in Kubernetes.



- Examples

- Environment variables set up in spec for admin and root passwords
- Create secrets for both
  - Create via command line

```
k create secret generic mysqlsecret2 --  
from-literal=mysqlpassword=admin --  
from-literal=mysqlrootpassword=root+1  
-n roar
```

- Create via yaml file

```
1 apiVersion: v1  
2 kind: Secret  
3 metadata:  
4   name: mysqlsecret  
5 type: Opaque  
6 data:  
7   mysqlpassword: YWRtaW4=  
8   mysqlrootpassword: cm9vdCsx
```

- Update spec to use those

```
58  
59   env:  
60     - name: MYSQL_DATABASE  
61       value: registry  
62     - name: MYSQL_PASSWORD  
63       value: admin  
64     - name: MYSQL_ROOT_PASSWORD  
65       value: root+1  
66     - name: MYSQL_USER  
67       value: admin
```

```
58  
59   env:  
60     - name: MYSQL_DATABASE  
61       value: registry  
62     - name: MYSQL_PASSWORD  
63       valueFrom:  
64         secretKeyRef:  
65           name: mysqlsecret  
66           key: mysqlpassword  
67     - name: MYSQL_ROOT_PASSWORD  
68       valueFrom:  
69         secretKeyRef:  
70           name: mysqlsecret  
71           key: mysqlrootpassword  
72     - name: MYSQL_USER  
73       value: admin
```



# Configmaps

59

- Examples

- Environment variables set up in spec for database and user
- Create configmap for both
  - Create via command line

```
k create configmap mysql-configmap --from-literal=mysql.database=registry  
--from-literal=mysql.user=admin -n roar To change them, must modify  
container
```

OR

- Create via yaml file

```
1 apiVersion: v1  
2 kind: ConfigMap  
3 metadata:  
4   name: mysql-configmap  
5 data:  
6   mysql.database: registry  
7   mysql.user: admin
```

- Update spec to use those

```
58   env:  
59     - name: MYSQL_DATABASE  
60       value: registry  
61     - name: MYSQL_PASSWORD  
62       valueFrom:  
63         secretKeyRef:  
64           name: mysqlsecret  
65           key: mysqlpassword  
66     - name: MYSQL_ROOT_PASSWORD  
67       valueFrom:  
68         secretKeyRef:  
69           name: mysqlsecret  
70           key: mysqlrootpassword  
71     - name: MYSQL_USER  
72       value: admin
```

```
58   env:  
59     - name: MYSQL_DATABASE  
60       valueFrom:  
61         configMapKeyRef:  
62           name: mysql-configmap  
63           key: mysql.database  
64     - name: MYSQL_PASSWORD  
65       valueFrom:  
66         secretKeyRef:  
67           name: mysqlsecret  
68           key: mysqlpassword  
69     - name: MYSQL_ROOT_PASSWORD  
70       valueFrom:  
71         secretKeyRef:  
72           name: mysqlsecret  
73           key: mysqlrootpassword  
74     - name: MYSQL_USER  
75       valueFrom:  
76         configMapKeyRef:  
77           name: mysql-configmap  
78           key: mysql.user
```



## Lab 3 – Working with Kubernetes secrets and configmaps

Purpose: In this lab, we'll get some practice storing secure and insecure information in a way that is accessible to k8s but not stored in the usual deployment files.



# Kubernetes Storage Concepts

- Container filesystem
  - Storage at the level of a container
    - » Ephemeral – only exists while container is running
    - » Any changes made are gone when container is gone
- Volume
  - Storage at the level of a pod
    - » Data is preserved if a container crashes
    - » Allows sharing of data across containers in a pod
    - » Only exists while pod exists
    - » Any changes made are gone when pod is gone (deleting pod deletes the volume)
- Persistent Volume
  - Storage outside the pod – cluster level
    - » Data is preserved if pod goes away
- Persistent Volume Claim
  - Way of saying what is needed for storage
    - » Defines specific amount of storage requested, access mode, etc.
  - Kubernetes looks for a specific matching Persistent Volume and matches it to claim
  - PVC's have different access modes they can specify
    - » ReadWriteOnce – one bound Pod with R/W access
    - » ReadOnlyMany – many bound Pods with R only access
    - » ReadWriteMany – many bound Pods with R/W access

```

1  ---
2  kind: PersistentVolume
3  apiVersion: v1
4  metadata:
5    name: mysql-pv
6    labels:
7      type: local
8  spec:
9    storageClassName: manual
10   capacity:
11     storage: 100M
12   accessModes:
13     - ReadWriteOnce
14   hostPath:
15     path: "/mnt/data"
16 ---
17  apiVersion: v1
18  kind: PersistentVolumeClaim
19  metadata:
20    labels:
21      app: mysql
22      name: mysql-pv-claim
23  spec:
24    storageClassName: manual
25    accessModes:
26      - ReadWriteOnce
27    resources:
28      requests:
29        storage: 100M
30
31  spec:
32    containers:
33      - name: mysql
34        image: localhost:5000/roar-db:v1
35    ports:
36      - name: mysql
37        containerPort: 3306
38    env:
39      - name: MYSQL_DATABASE
40        valueFrom:
41          configMapKeyRef:
42            name: mysql-configmap
43            key: mysql.database
44      - name: MYSQL_PASSWORD
45        valueFrom:
46          secretKeyRef:
47            name: mysqlsecret
48            key: mysqlpassword
49      - name: MYSQL_ROOT_PASSWORD
50        valueFrom:
51          secretKeyRef:
52            name: mysqlsecret
53            key: mysqlrootpassword
54      - name: MYSQL_USER
55        valueFrom:
56          configMapKeyRef:
57            name: mysql-configmap
58            key: mysql.user
59    volumeMounts:
60      - mountPath: /var/lib/mysql
61        name: mysql-pv-claim
62    volumes:
63      - name: mysql-pv-claim
64    persistentVolumeClaim:
65      claimName: mysql-pv-claim
66
67  ---

```



- Objects
  - PV's
    - physical volume on the host machine or network that stores the data to be persisted
  - PVC's
    - request for PV's
    - claim check for resource
  - PV's are attached to pods via PVC's
  - Storageclasses define the type of PVCs that a user can request
  - Think of it like:

**Pod -> PVC -> PV -> (Storage Class, If applicable) -> Host machine**



# Storage Classes and Provisioners

63

- Storage classes provides way for provider to define classes of storage they offer
- May define QOS levels, backup policies, etc.
- Storage classes have a Provisioner to determine what kind of PV
- Storage classes can dynamically provision a PV if so configured
- Admins can specify StorageClass for PVCs

Volume Plugin	Internal Provisioner	Config Example
AWSBlockStore	✓	AWS EBS
AzureFile	✓	Azure File
AzureDisk	✓	Azure Disk
CephFS	-	-
Cinder	✓	OpenStack Cinder
FC	-	-
Flexvolume	-	-
Flocker	✓	-
GCEPersistentDisk	✓	GCE PD
Glusterfs	✓	Glusterfs
iSCSI	-	-
Quobyte	✓	Quobyte
NFS	-	-
RBD	✓	Ceph RBD
VsphereVolume	✓	vSphere
PortworxVolume	✓	Portworx Volume
ScaleIO	✓	ScaleIO
StorageOS	✓	StorageOS
Local	-	Local

## Lab 4 – Working with persistent storage - Kubernetes Persistent Volumes and Persistent Volume Claims

Purpose: In this lab, we'll see how to connect pods with external storage resources via persistent volumes and persistent volume claims.



# Helm



# What is Helm?

66

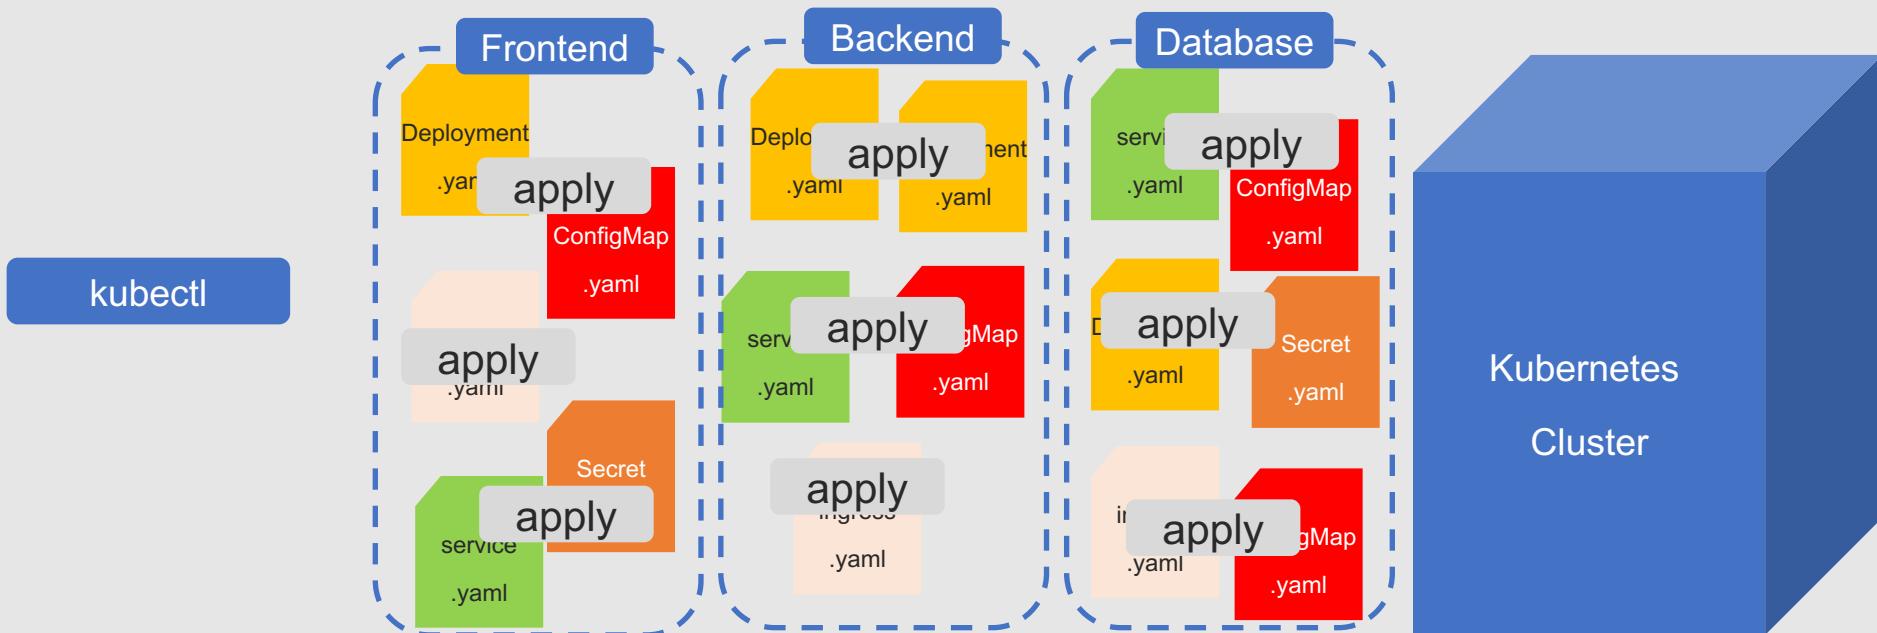
- Package Manager and Lifecycle Manager for K8s
  - Like yum, apt but for K8s
  - Bundles related manifests (such as deployment.yaml, service.yaml, etc.) into a “chart”
  - When installing chart, Helm creates a “release”
  - Lifecycle management
    - Create, Install, Upgrade, Rollback, Delete, Status, Versioning
  - Benefits
    - Templating, Repeatability, Reliability, Multiple Environment, Ease of collaboration



# Why do we need something like this?

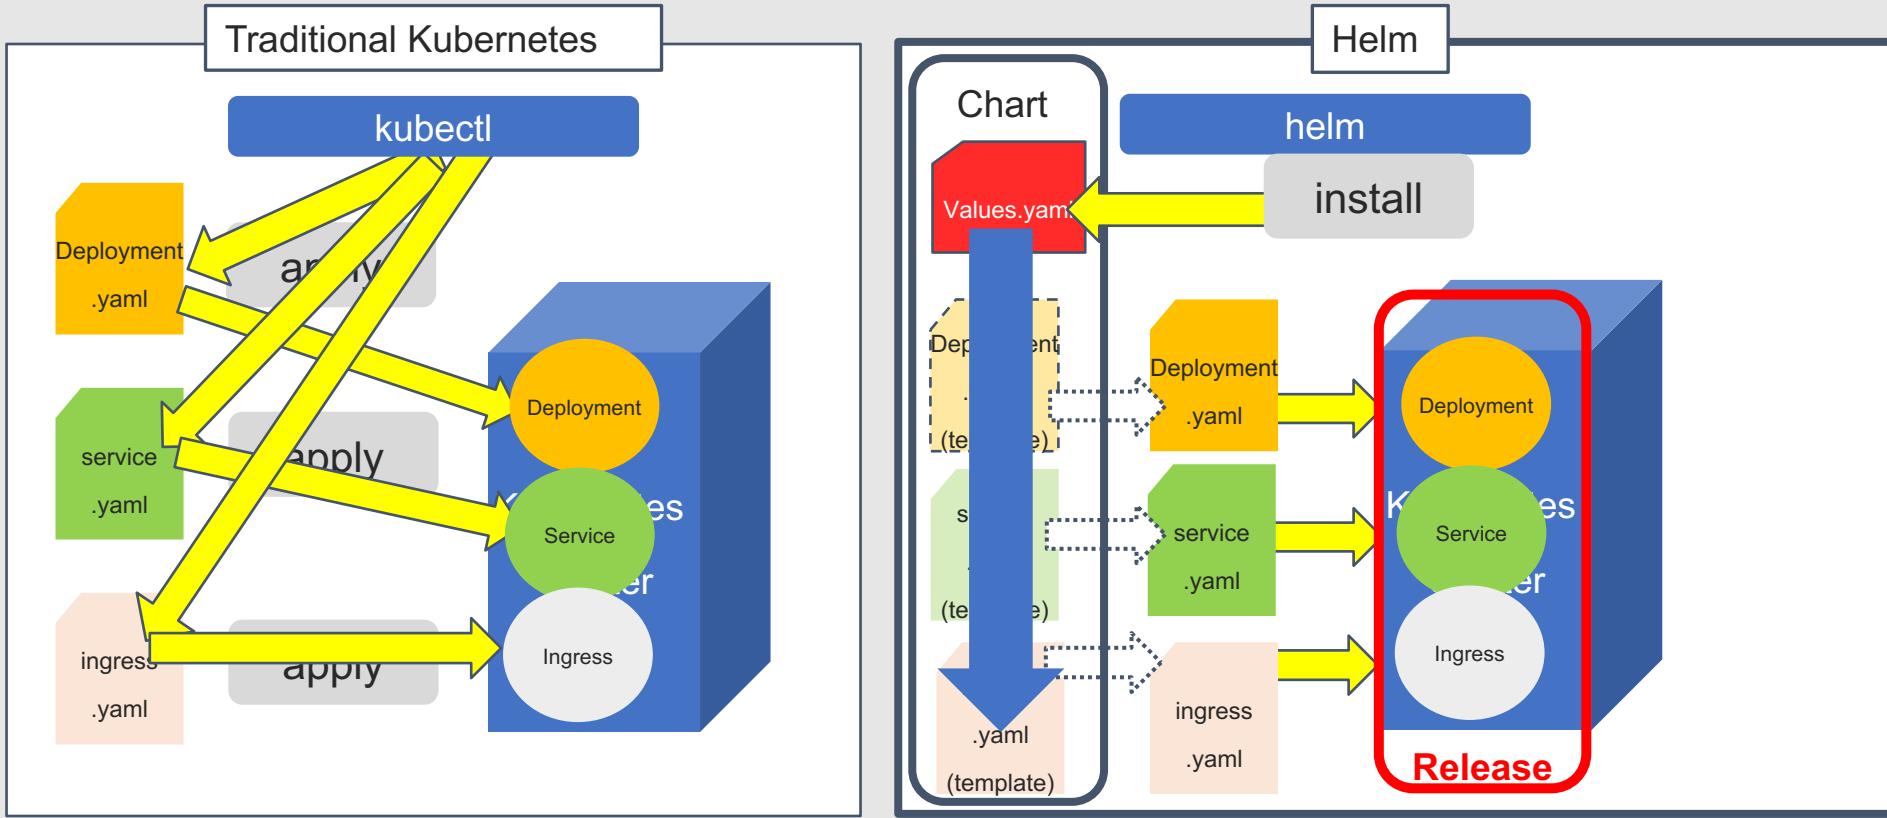
67

- Scale and complexity
- Microservice = Pod + Deployment+ ReplicationSet + Ingress + Service times # of microservices
- Duplication of values across objects
- Hard to override values (no parameterization)
- Managing lifecycle of all the objects is challenging



# How does Helm simplify things?

- Traditional deployment in Kubernetes is done with kubectl across files into separately managed items
- Helm deploys units called charts as managed releases





# What are the advantages of using Helm?

69

- Saves having to deploy multiple individual manifests (for multiple Kubernetes objects)
- Allows for reuse via parameterizing (templates)
- Manages releases of Helm packages
- Simplified mechanism for finding and deploying popular software (packaged as Helm charts)
- Makes it easier to share your applications (via charts)



# Helm Topology

70

- Chart – a package; a bundle of K8s resources
- Release – a chart instance loaded into K8s
  - Same chart can be installed several times into the same cluster
  - Each such chart will have its own release
- Repository – a repository of published charts
- Template - a K8s configuration file mixed with Go/Spring templates

NAME	REVISION	UPDATED	STATUS	CHART	APP VERSION	NAMESPACE
istio	2	Sat Jul 20 22:09:38 2019	DEPLOYED	istio-1.2.0	1.2.0	istio-system
istio-init	1	Thu Jun 27 13:34:49 2019	DEPLOYED	istio-init-1.2.0	1.2.0	istio-system
istiol	5	Sun Oct 27 17:56:53 2019	DEPLOYED	roar-web-0.1.0		istiol
jenkins-x	1	Thu Jun 6 07:53:23 2019	DEPLOYED	jenkins-x-platform-2.0.330		jx
roar2	2	Sun Oct 27 17:31:35 2019	DEPLOYED	roar-helm-0.1.0		roar2





# What is a Chart?

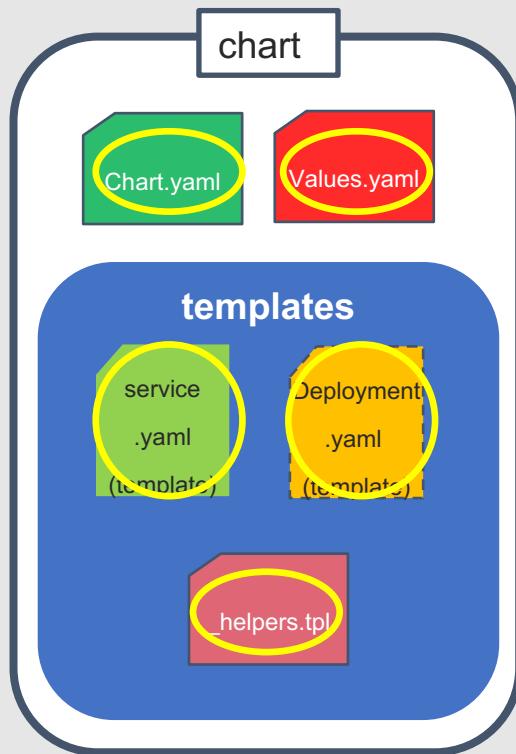
71

- Packaging format for Helm
- Define a way to compose a set of K8S resources and values to make up a deployment
- Deployable unit
  - Can be installed, updated, removed
- “Source code”
  - Can be versioned and managed in source control and compressed in packages
- Charts can include other charts as dependencies



# Example Files in a Helm Chart

- Main chart directory
- Chart.yaml – chart description
- Values.yaml – values to be used in chart
- Templates – templated files that will form K8s manifests
- \_helpers.tpl – helper functions



**deployment.yaml**

```

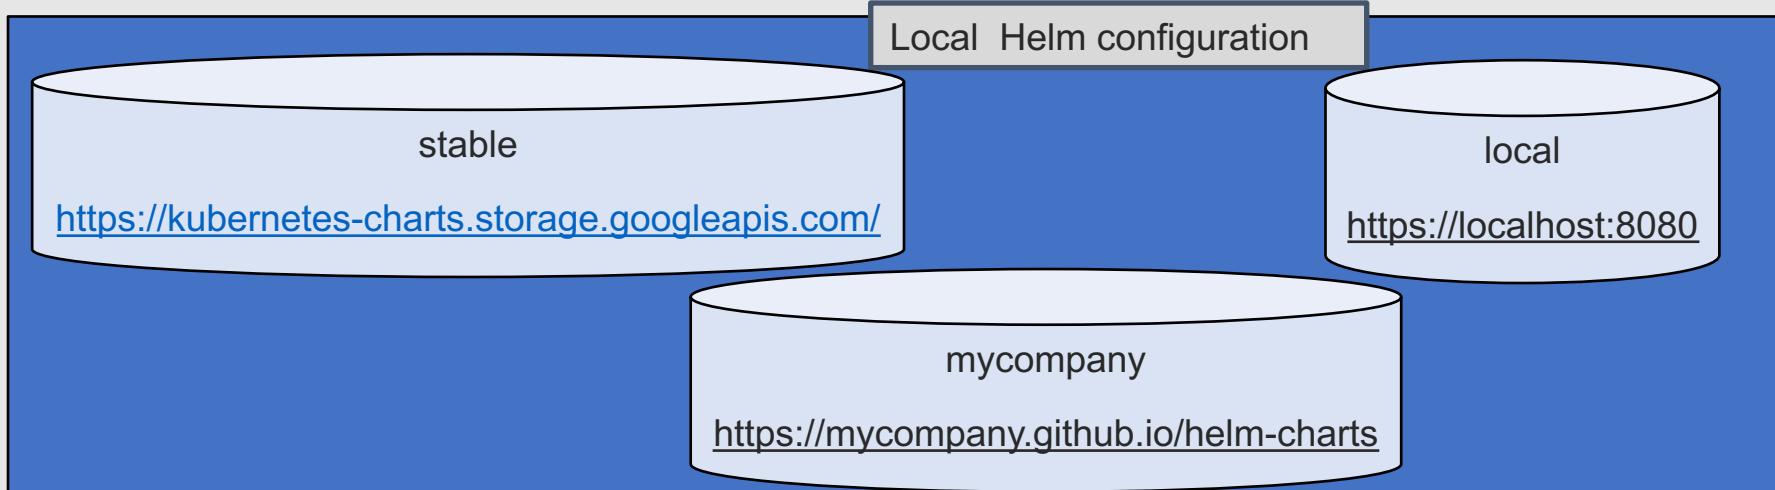
1 apiVersion: extensions/v1beta1
2 kind: Deployment
3 metadata:
4   name: {{ template "roar-db.name" . }}
5   labels:
6     app: {{ template "roar-db.name" . }}
7     chart: {{ .Chart.Name }}-{{ .Chart.Version | replace "+" "_" }}
8     release: {{ .Release.Name }}
9     namespace: {{ .Values.namespace }}
10 spec:
11   replicas: {{ .Values.replicaCount }}
12   template:
13     metadata:
14       labels:
15         app: {{ template "roar-db.name" . }}
16   spec:
17     containers:
18       - name: {{ .Chart.Name }}
19         image: bclaster/roar-db-image:v1
20         imagePullPolicy: Always
21       ports:
22         - name: {{ .Values.deployment.ports.name }}
23           containerPort: {{ .Values.deployment.ports.containerPort }}
24   env:
25     {{- include "roar-db.environment-values" . | indent 10 }}

```



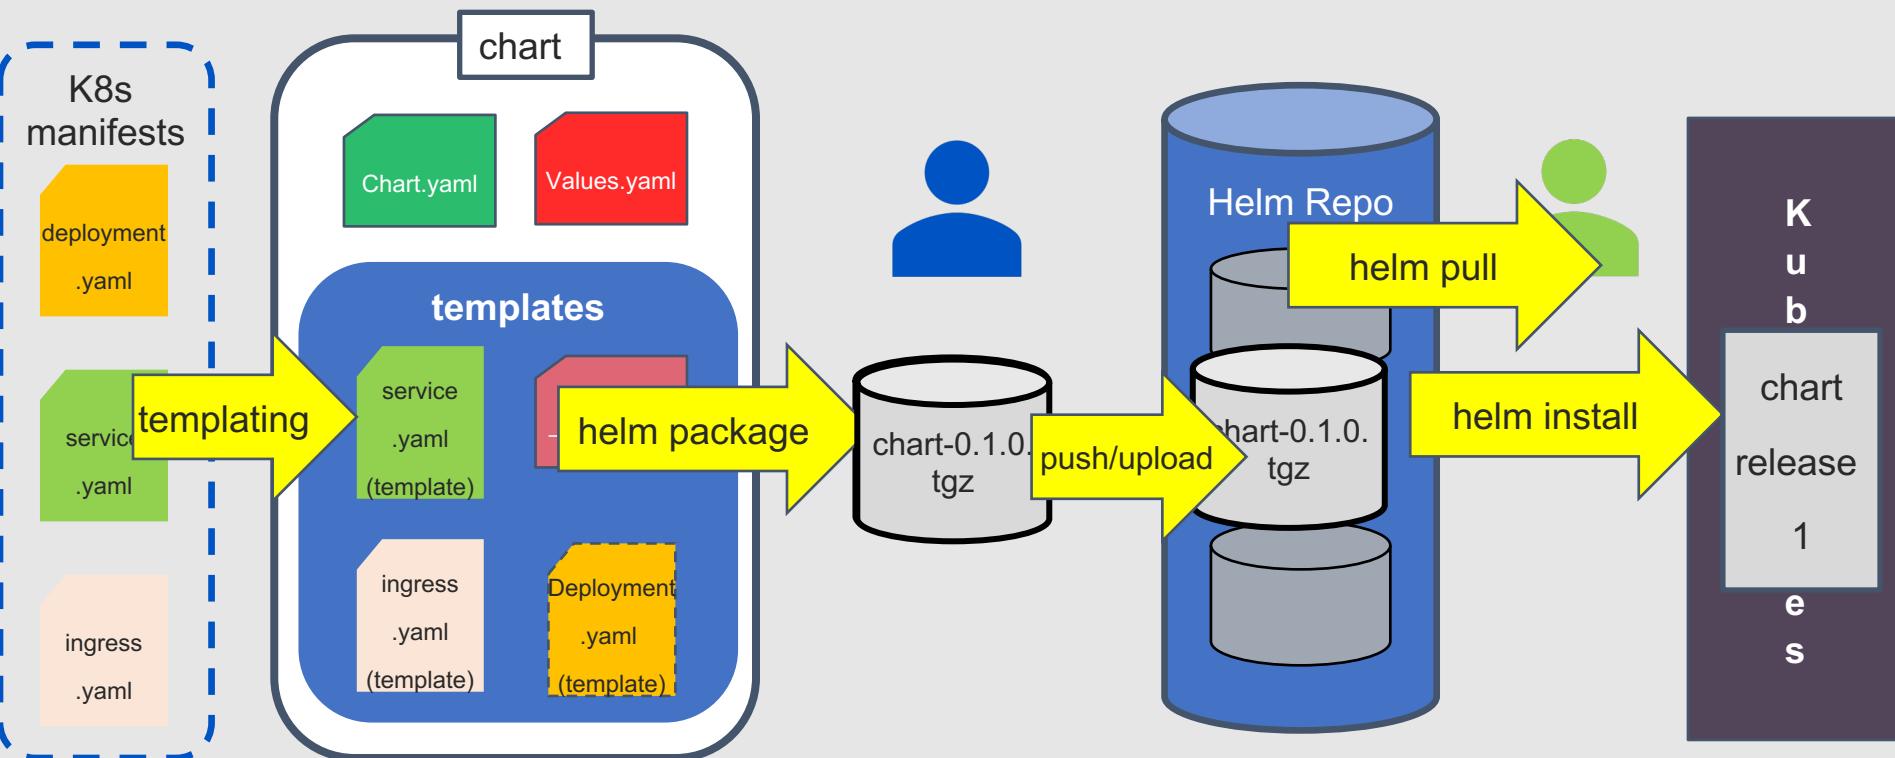
# What is a chart repository?

- Chart repo
  - Location where Helm charts can be stored and shared
  - Able to serve yaml and tar files
    - » Helm package format is tar file
    - » Index file for repo is yaml
  - Responds to REST API GET requests to get packages
  - Many storage options available – cloud buckets, local storage, GH Pages, etc.
- Helm instance can have many repos defined/added



# Helm as Package Manager

- Kubernetes files are “templated”
- Templates and related files are structured as a chart
- Charts are packaged
- Packages are stored in repos for easy use by others
- Charts can be pulled (downloaded) and optionally unpacked (untar)
- Charts are installed from repos (with values) as releases into Kubernetes





# Helm Repo Operations

- Show all repos
  - \$ helm repo list
- Add a repo
  - \$ helm repo add <repo> url
  - \$ helm repo add stable <https://kubernetes-charts.storage.googleapis.com/>
- Search a repo
  - \$ helm search repo <repo> <chartname>
  - \$ helm search hub
- Remove a repo
  - \$ helm repo rm <repo>

```
$ helm search repo stable | head
```

NAME	CHART VERSION	APP VERSION	DESCRIPTION
stable/acs-engine-autoscaler	2.2.2	2.1.1	DEPRECATED Scales worker nodes within agent pools
stable/aerospike	0.3.2	v4.5.0.5	A Helm chart for Aerospike in Kubernetes
stable/airflow	7.1.5	1.10.10	Airflow is a platform to programmatically autho...
stable/ambassador	5.3.2	0.86.1	DEPRECATED A Helm chart for Datawire Ambassador
stable/anchore-engine	1.6.9	0.7.2	Anchore container analysis and policy evaluatio...
stable/apm-server	2.1.5	7.0.0	The server receives data from the Elastic APM a...
stable/ark	4.2.2	0.10.2	DEPRECATED A Helm chart for ark



# Helm Operations

- completion generate autocompletions script for the specified shell (bash or zsh)
- create create a new chart with the given name
- dependency manage a chart's dependencies
- env helm client environment information
- get download extended information of a named release
- help Help about any command
- history fetch release history
- **install install a chart**
- lint examine a chart for possible issues
- **list list releases**
- package package a chart directory into a chart archive
- plugin install, list, or uninstall Helm plugins
- pull download a chart from a repository and (optionally) unpack it in local directory
- **repo add, list, remove, update, and index chart repositories**
- rollback roll back a release to a previous revision
- **search search for a keyword in charts**
- **show show information of a chart**
- status display the status of the named release
- template locally render templates
- test run tests for a release
- uninstall uninstall a release

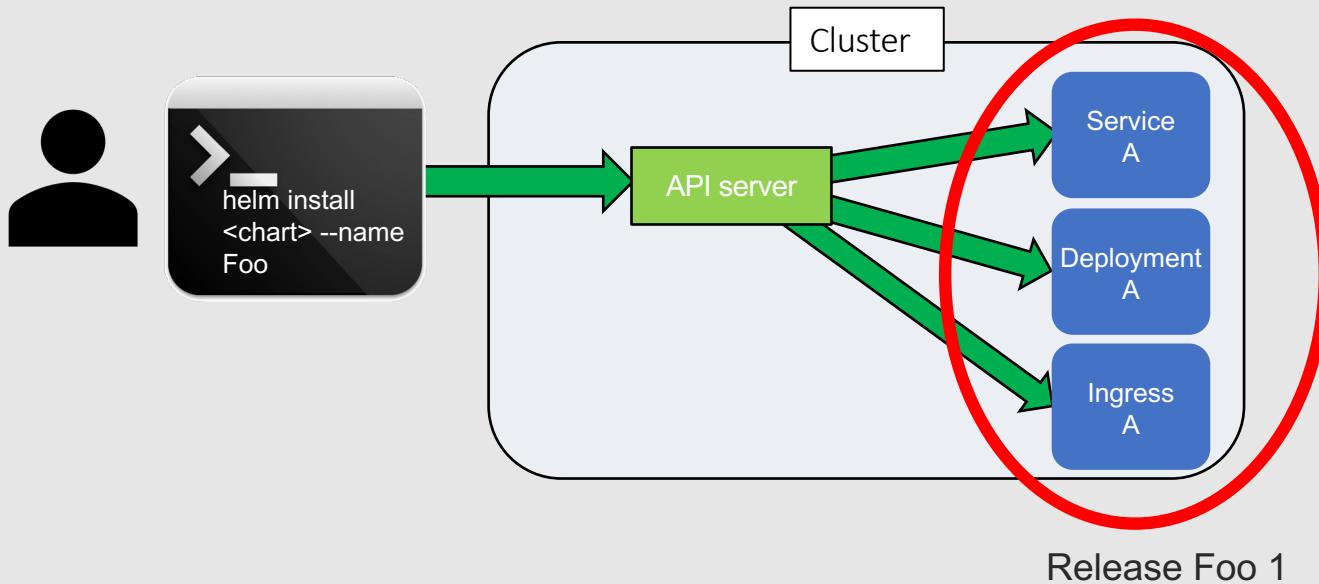




# Helm Install a Chart

77

- helm install <chart>





# Getting Information about Releases

78

- A release is an instance of a Helm chart deployed in Kubernetes
- Has a release name that can be different from chart name
- See list of releases
  - \$ helm list
- See current status
  - \$ helm status RELEASE\_NAME [flags]
- See history of release
  - \$ helm history RELEASE\_NAME [flags]

```
$ helm status local-chartmuseum
NAME: local-chartmuseum
LAST DEPLOYED: Sun Jun 14 15:55:32
2020
NAMESPACE: default
STATUS: deployed
REVISION: 2
TEST SUITE: None
NOTES:
** Please be patient while the chart is
being deployed **
```

Get the ChartMuseum URL by running:

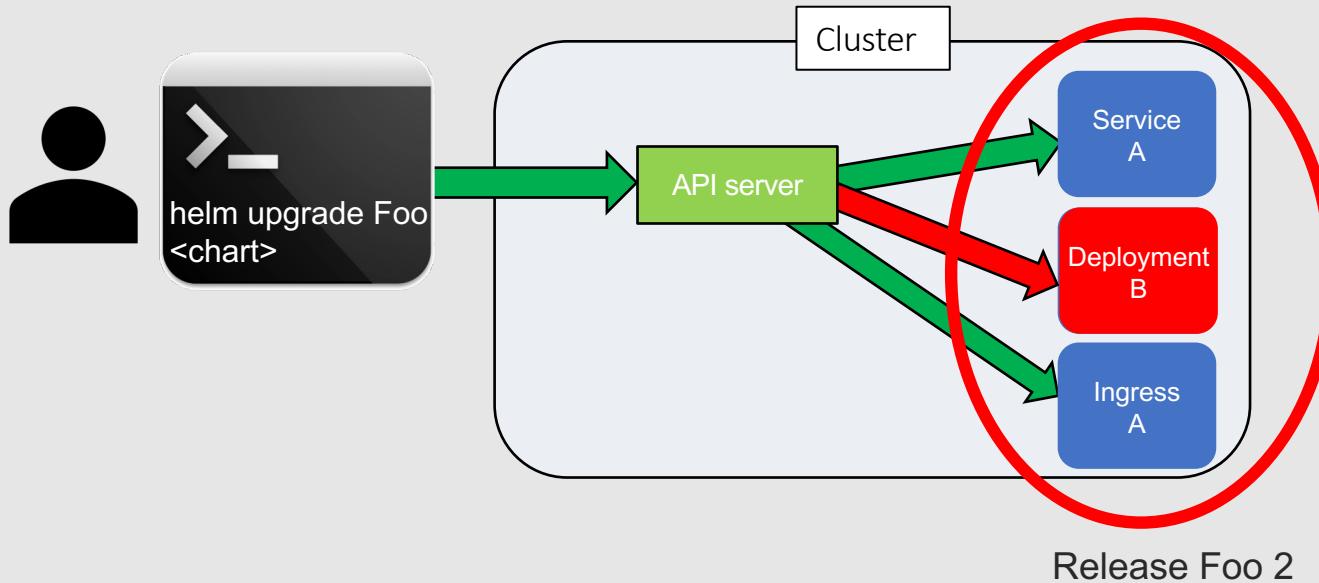
\$ helm history local-chartmuseum						
REVISION	UPDATED	STATUS	CHART	APP VERSION	DESCRIPTION	
1	Sun Jun 14 15:45:38 2020	superseded	chartmuseum	2.13.0	0.12.0	Install complete
2	Sun Jun 14 15:55:32 2020	deployed	chartmuseum	2.13.0	0.12.0	Upgrade complete



# Helm Upgrading a Chart

79

- `helm upgrade <release> <chart>`
- use `--set` to override chart settings

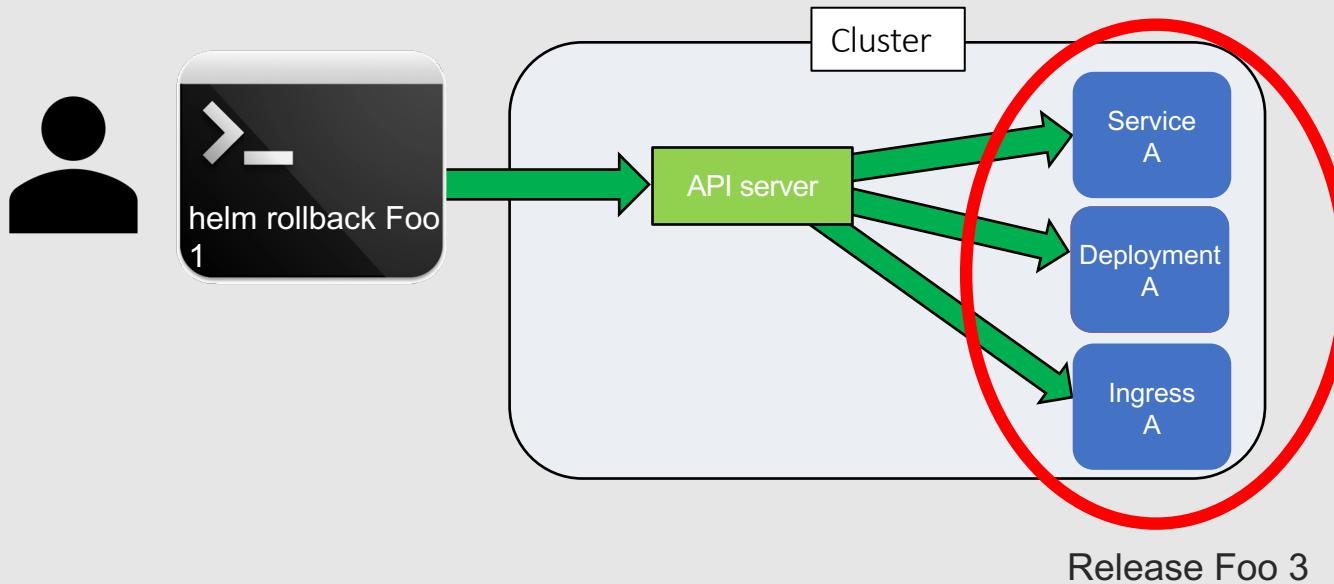




# Helm Rollback

80

- `helm history <release>` (to see revision numbers)
- `helm rollback <release> <revision>`





# Our Example App

- R.O.A.R. (Registry of Animal Responders)
- Simple app with two pieces
  - Webapp on the front
  - Mysql database on the backend
- Written in Java
- Simple REST API
- War file is produced for webapp
- Managed in Tomcat

ROAR Agents

172.17.0.17:8080/roar/

R.O.A.R (Registry of Animal Responders) Agents

Show 10 entries Search:

ID	Name	Species	Date of First Service	Date of Last Service	Adversary	Adversary Tech
1	Road Runner	bird	1955-01-20	1995-02-15	Wile E. Coyote	ACME product du jour
2	Scooby	dog	1969-05-19	2000-02-11	fake ghosts	mask
3	Perry	platypus	2013-01-20	2015-04-09	H. Doofensmirtz	...inator
4	Mr. Krabs	crab	2010-06-17	2014-07-07	Plankton	various
5	Bugs Bunny	rabbit	1966-05-22	1988-04-15	E. Fudd	wabbit gun

Showing 1 to 5 of 5 entries Previous 1 Next

## Lab 5 - Working with Helm

Purpose: In this lab, we'll compare a Helm chart against standard Kubernetes manifests and then deploy the Helm chart into Kubernetes



# Traditional K8s yaml vs Helm template

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: roar-web
  labels:
    app: roar-web
  namespace: roar
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: roar-web
    spec:
      containers:
        - name: roar-web
          image: localhost:5000/roar-web-v1
        ports:
          - name: web
            containerPort: 8080
```

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: {{ template "roar-web.name" . }}
  labels:
    app: {{ template "roar-web.name" . }}
  chart: {{ .Chart.Name }}-{{ .Chart.Version | replace "+" "_" }}
  release: {{ .Release.Name }}
  namespace: {{ .Values.namespace }}
spec:
  replicas: {{ .Values.replicaCount }}
  template:
    metadata:
      labels:
        app: {{ template "roar-web.name" . }}
    spec:
      containers:
        - name: {{ .Chart.Name }}
          image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
          {{- if .Values.image.pullPolicy }}
            imagePullPolicy: {{ toYaml .Values.image.pullPolicy }}
          {{- end }}
          ports:
            - name: {{ .Values.deployment.ports.name }}
              containerPort: {{ .Values.deployment.ports.containerPort }}
```

# How Values are resolved in Helm

spec:

containers:

- name: {{ .Chart.Name }}

image: bclaster/roar-db-image:v1

imagePullPolicy: Always

ports:

- name: {{ .Values.deployment.ports.name }}

containerPort: {{ .Values.deployment.ports.containerPort }}

env:

{-{ include "roar-db.environment-values" indent 10 }}

values.yaml

# Default values for roar-db chart.

# This is a YAML-formatted file.

# Declare variables to be passed

# into your templates.

replicaCount: 1

nameOverride: mysql

deployment:

ports:

name: mysql

containerPort: 3306

templates/deployment.yaml

- string indicates hierachial path
- dot notation separates levels
- .Values refers to top-level values from values.yaml

deployment.yaml

...

containerPort: 3306



# Template Functions

85

- Allows you to transform data passed into a template
- Function syntax
  - `functionname arg1 arg2 ...`
- Over sixty functions available to use; provided by
  - Go template language
  - Go sprig template library
- Examples
  - `upper`, `quote`, `eq`
  - custom functions



# Pipelines

- Similar to the idea of Unix pipelines
- Allows chaining together multiple template commands
- Simplify doing multiple things in sequence
- When pipelining arguments, result of previous evaluation is sent as last argument to next one

`{{ first evaluation | function arg1 }}`

equivalent to

`{{ function arg1 arg2 }}` where “arg2” = “first evaluation”

- Example

name: `{{ .Values.user.name | upper | quote }}`

## Lab 6 - Templating with Helm

Purpose: In this lab, you'll get to see how we can change hard-coded values into templates, override values, and upgrade releases through Helm.



# Kustomize



# Reuse from Customization

89

- Kubernetes yaml files are often reused
- Approaches for reuse
  - Modify the original to suit your needs
  - Make multiple copies and edit them separately for different needs
- Challenges
  - Divergence from original “source” files
  - Having multiple disparate copies means it is more challenging to leverage a common “base” set
  - Making consistent changes across multiple, independent copies for new versions, etc. becomes challenging
- Another approach is parameterizing templates – i.e. Helm



# Challenges with Helm

90

- Templates and value files are not usable Kubernetes specifications
- They are more like a new kind of “programming” that wraps around the Kubernetes spec
- Parameterized templates change the original specification
- Many things become parameters for many different users
- Counters a typical goal of reuse – keep differences between specs small and simple

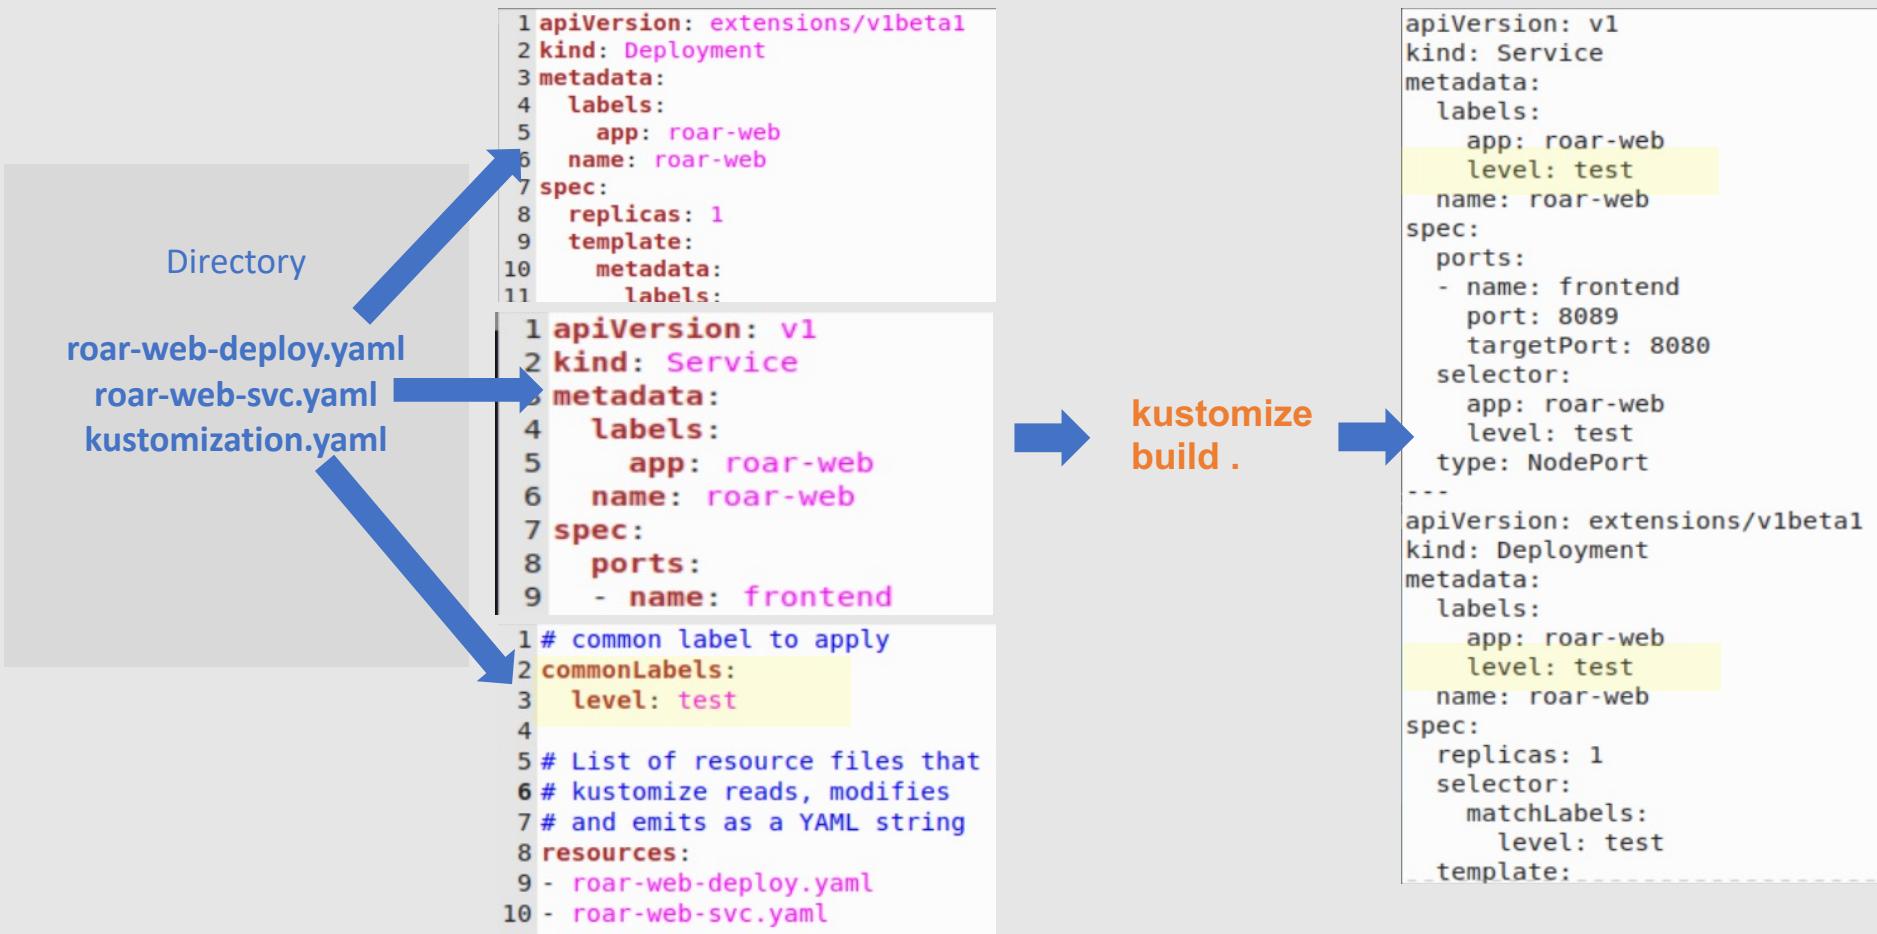


# Another Option - Kustomize

- Kustomize creates modified Kubernetes manifests by “overlaying” declarative specifications on top of existing Kubernetes manifests
- Changes to make/overlay are stored in a separate file called *kustomization.yaml*
- Kustomize reads the *kustomization.yaml* file and the Kubernetes spec files (manifests) and dumps out completed resources to stdout with the changes from *kustomization.yaml* applied “on top of” the standard Kubernetes files
- *kustomize build* command causes files to be processed and output created
- Kustomize leverages the idea of “transformers”
- Transformers are functionality built into Kustomize that “transforms” your Kubernetes manifests given a simple set of declarative rules



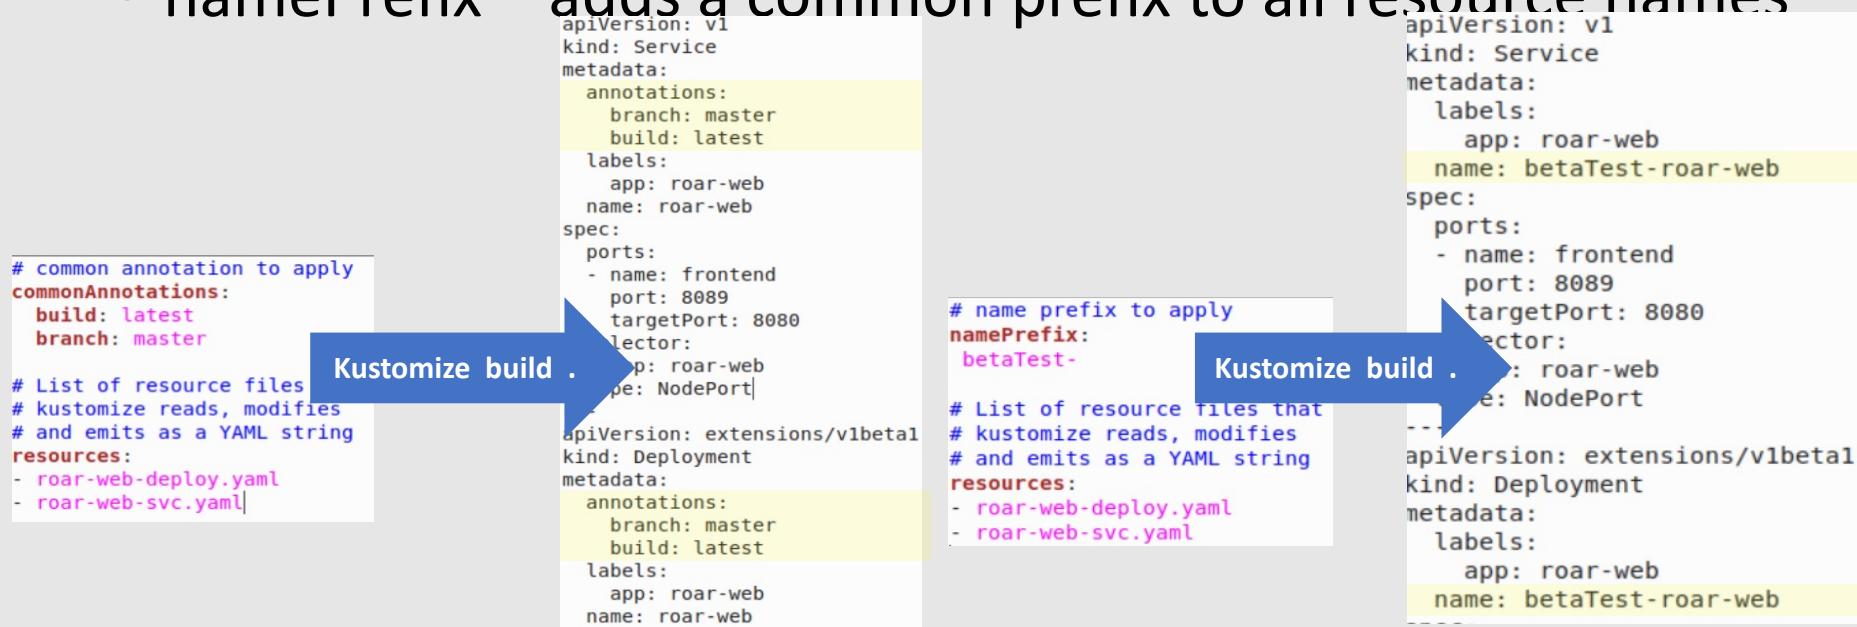
# Simple Example





# "Common" changes

- commonLabel – adds a common label (name: value) to each K8S resource
- commonAnnotations – adds an annotation to all K8S resources
- namePrefix – adds a common prefix to all resource names



## Lab 7 - Run a basic Kustomize example

Purpose: In this lab, we'll see how to make a set of manifests usable with Kustomize and how to use Kustomize to add additional changes without modifying the original files.



- Common use case is needing multiple variations (variants) of a set of common resources (i.e. dev, stage, prod)
- To do this, Kustomize has concept of “overlay” and “base”
  - Both represented via kustomization.yaml file (includes resources + customizations)
  - Base – declares things variants have in common
  - Overlays – declare the differences
- Variants can also apply “patches”
  - Patch – partial deployment spec – used to update some existing spec as opposed to adding new information



# Kustomize Base

```
diyuser3@training1:~/ato-ws/roar-kz/base/db$ kustomize build .
```

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: roar-db
  1  apiVersion: v1
  2  kind: Service
  3  metadata:
  4  labels:
  5    app: roar-db
  6  name: roar-db
  7 spec:
  8   ports:
  9     - name: mysql
 10    port: 3306
 11    targetPort: 3306
 12 selector:
 13   app: roar-db
```

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4  labels:
5    app: roar-db
6    name: roar-db
7  spec:
8    ports:
9      - name: mysql
10     port: 3306
11     targetPort: 3306
12   app: roar-db
13   path: "/metadata/name"
14   value: "mysql"
15
16
17  name: "roar-db"
```

hat kustomize reads, modifies  
ng

Resources to  
include

from "roar-db" to "mysql"  
ted with the webapp

Patch block = defines  
Type of patch  
patch in existing text  
hierarchy and  
replacement value

Item to search for



# Kustomize Base

- Original resources + kustomization

diyuser3@training1:~/ato-ws/roar-kz/base\$ kustomize build .

```
1 # Common label to be applied
2 commonLabels:
3   app: original
4
5 apiVersion: v1
6 kind: Namespace
7 metadata:
8   #   labels:
9     app: original
10    #   name: roar-original
11
12 ---
```

modifies

```
13   apiVersion: v1
14   kind: Service
15   metadata:
16     labels:
17       app: original
18       name: mysql
19       namespace: roar-original
```



# Kustomize Base

```
diyuser3@training1:~/ato-ws/roar-kz/base/web$ kustomize build . . . . .  
apiVersion: v1  
kind: Service  
metadata:  
  labels:  
    app: roar-web  
  name: roar-web  
spec:  
  ports:  
  - name: frontend  
    port: 8089  
    targetPort: 8080  
  selector:  
    app: roar-web  
  type: NodePort  
---  
apiVersion: extensions/v1beta1  
kind: Deployment  
metadata:  
  labels:  
    app: roar-web  
  name: roar-web  
spec:  
  replicas: 1  
  template:  
    metadata:  
      labels:  
        app: roar-web  
    spec:  
      containers:  
      - image: localhost:5000/roar-web:v1  
        imagePullPolicy: Always  
        name: roar-web  
      ports:  
      - containerPort: 8080  
        name: web  
  
piVersion: extensions/v1beta1  
---  
apiVersion: v1  
kind: Service  
metadata:  
  labels:  
    app: roar-web  
  name: roar-web  
1  
2  # List of resource files that kustomize reads, modifies  
3  # and emits as a YAML string  
4  resources:  
5  - roar-web-deploy.yaml  
6  - roar-web-svc.yaml  
  
selector:  
  app: roar-web  
  type: NodePort  
  - containerPort: 8080  
  name: web  
  
Resources to include
```



# Kustomize Feature List

Field	Type	Explanation
namespace	string	add namespace to all resources
namePrefix	string	value of this field is prepended to the names of all resources
nameSuffix	string	value of this field is appended to the names of all resources
commonLabels	map[string]string	labels to add to all resources and selectors
commonAnnotations	map[string]string	annotations to add to all resources
resources	[]string	each entry in this list must resolve to an existing resource cfg file
configmapGenerator	[] <a href="#">ConfigMapArgs</a>	Each entry in this list generates a ConfigMap
secretGenerator	[] <a href="#">SecretArgs</a>	Each entry in this list generates a Secret
generatorOptions	<a href="#">GeneratorOptions</a>	Modify behaviors of all ConfigMap and Secret generator
bases	[]string	Each entry in this list should resolve to a directory containing a kustomization.yaml file
patchesStrategicMerge	[]string	Each entry in this list should resolve a strategic merge patch of a Kubernetes object
patchesJson6902	[] <a href="#">Json6902</a>	Each entry in list should resolve to a Kubernetes obj and a Json Patch
vars	[] <a href="#">Var</a>	Each entry is to capture text from one resource's field
images	[] <a href="#">Image</a>	Each entry is to modify the name, tags and/or digest for one image without creating patches
configurations	[]string	Each entry in this list should resolve to a file containing <a href="#">Kustomize transformer configurations</a>
crds	[]string	Each entry in this list should resolve to an OpenAPI definition file for Kubernetes types



# Kustomize Option in Kubernetes

- Kubectl has an option to reference Kustomize resources
  - “-k” or “--kustomize”
- Option needs to point to a kustomization directory
  - example: kubectl apply -k <kustomization directory>
- Multiple operations understand –k option
  - kubectl apply -k ./
  - kubectl get -k ./
  - kubectl describe -k ./
  - kubectl diff -k ./
  - kubectl delete -k ./

## Lab 8 - Creating Variants

Purpose: In this lab, we'll see how to create production and stage variants of our simple application

# Monitoring





# Kubernetes Dashboard

103

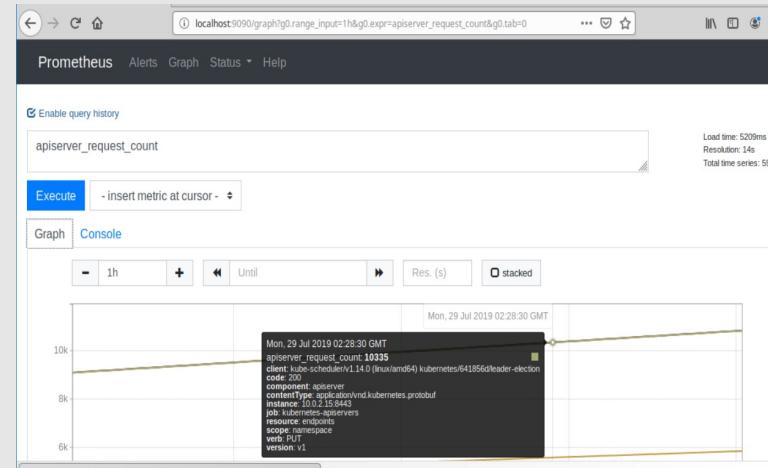
- Visual interface to the cluster
- General-purpose, web-based UI for clusters
- Can get overall status about objects in the cluster
- Invoke via “minikube dashboard” or “kubectl apply” from location and then “kubectl proxy”
- Can also create or modify objects via “+ CREATE” link
  - Can type in json or yaml
  - Can upload a spec file

The screenshot shows the Kubernetes Dashboard's Overview page. The left sidebar has a navigation menu with items: Cluster, Namespaces, Nodes, Persistent Volumes (which is selected), Roles, Storage Classes, Namespace (with 'default' selected), Workloads, and Cron Jobs. The main content area is divided into two sections: 'Discovery and Load Balancing' and 'Config and Storage'. The 'Discovery and Load Balancing' section contains a 'Services' table with one entry: 'kubernetes' (Name), component: kube (Labels), 10.96.0.1 (Cluster IP), kubernetes:44... (Internal endpoints), - (External endpoints), and a month (Age). The 'Config and Storage' section contains a 'Persistent Volume Claims' table with no visible entries.



- Event Monitoring and Alerting

- Open-source application for event monitoring and alerting
  - Records realtime metrics in a time-series database
  - Powerful query language - PromQL
  - Prometheus “monitoring platform” usually has
    - Multiple exporters that typically run on the monitored host to export metrics
    - Prometheus to centralize and store the metrics
    - Alertmanager to trigger alerts based on the metrics
    - Grafana to produce dashboards
- [https://en.wikipedia.org/wiki/Prometheus\\_\(software\)](https://en.wikipedia.org/wiki/Prometheus_(software))



MysqlHighSelectUsage (1 active)

```

name: MysqlHighSelectUsage
expr: rate(mysql_global_status_commands_total[commands=~"(select)"]{label} * 100 > 35
labels:
  severity: warning
annotations:
  description: MySQL High Level of SELECT usage {{ $labels.instance }}
  VALUE = {{ $value }}
  LABELS = {{ $labels }}
  summary: MySQL Select Usage High (instance {{ $labels.instance }})
  
```

Labels	State	Active Since	Value
alarmname=MysqlHighSelectUsage command=select instance=mysql-exporter-prometheus-mysql-exporter.monitoring.svc.cluster.local:9104	FIRING	2022-01-09T19:19:47.065791371Z	51.111111111111112

Annotations

```

description
MySQL High Level of SELECT usage mysql-exporter-prometheus-mysql-exporter.monitoring.svc.cluster.local:9104 VALUE = 51.111111111111112 LABELS = map[command:select instance:mysql-exporter-prometheus-mysql-exporter.monitoring.svc.cluster.local:9104 job:mysql]
summary
MySQL Select Usage High (instance mysql-exporter-prometheus-mysql-exporter.monitoring.svc.cluster.local:9104)
  
```



# Example Node Exporter Metrics

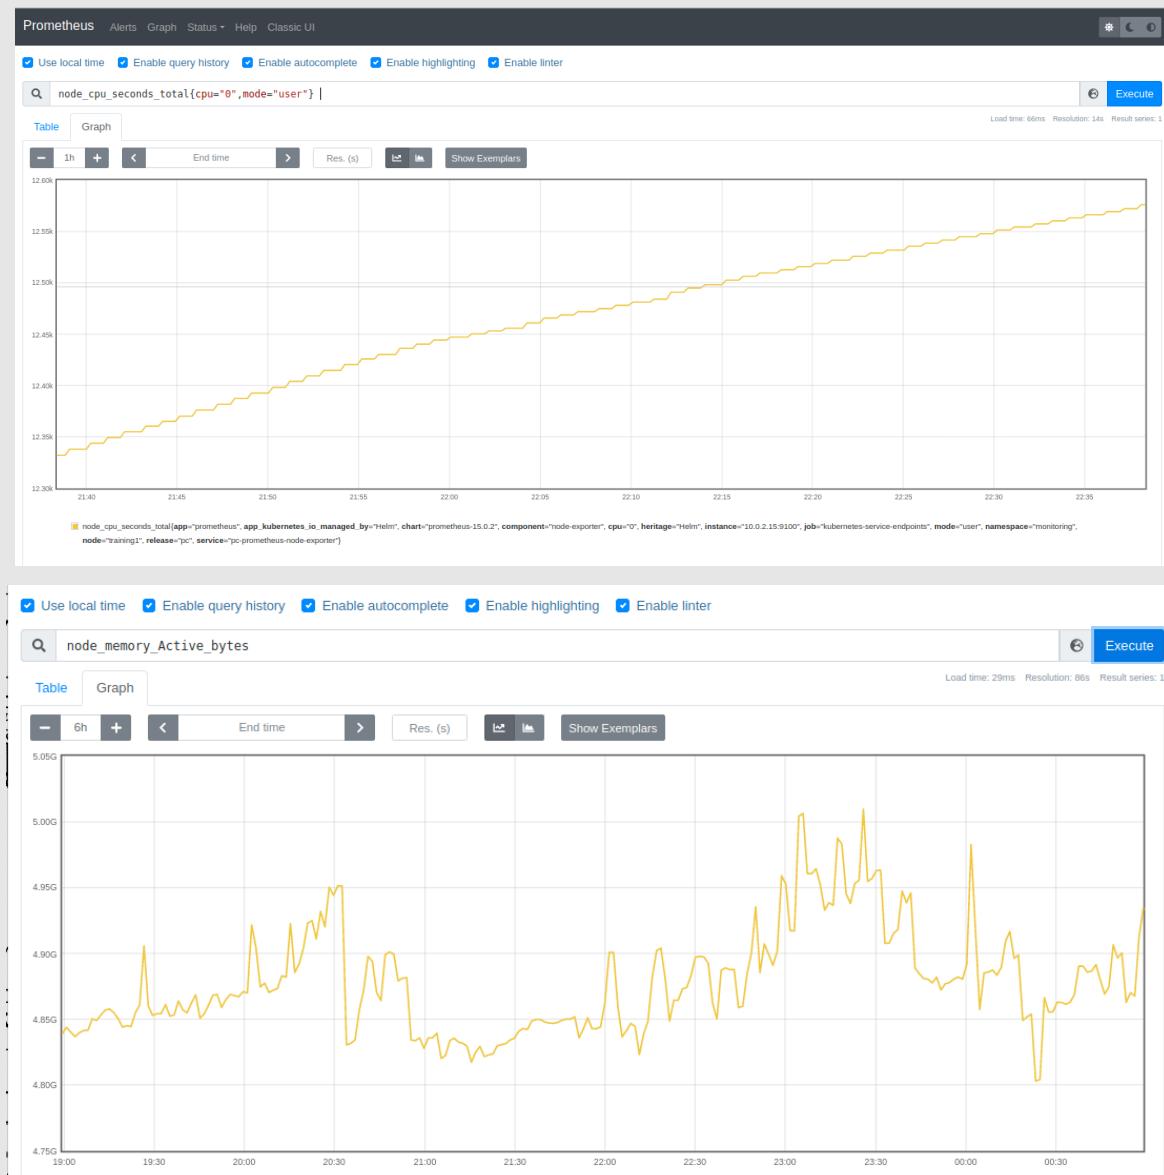
- Gauge
- Counter
- Multiple interfaces shown

```
localhost:9100/metrics
node_network_net_dev_group{device="veth40c6bad"} 0
node_network_net_dev_group{device="veth6999e49"} 0
node_network_net_dev_group{device="veth821ba89"} 0
node_network_net_dev_group{device="vetha68ccf1"} 0
node_network_net_dev_group{device="vetha9bee2f"} 0
node_network_net_dev_group{device="vethff841cb"} 0
# HELP node_network_protocol_type protocol_type value of /sys/class/net/<iface>.
# TYPE node_network_protocol_type gauge
node_network_protocol_type{device="docker0"} 1
node_network_protocol_type{device="enp0s3"} 1
node_network_protocol_type{device="lo"} 772
node_network_protocol_type{device="veth02b5642"} 1
node_network_protocol_type{device="veth2466b1a"} 1
node_network_protocol_type{device="veth40c6bad"} 1
node_network_protocol_type{device="veth6999e49"} 1
node_network_protocol_type{device="veth821ba89"} 1
node_network_protocol_type{device="vetha68ccf1"} 1
node_network_protocol_type{device="vetha9bee2f"} 1
node_network_protocol_type{device="vethff841cb"} 1
# HELP node_network_receive_bytes_total Network device statistic receive_bytes.
# TYPE node_network_receive_bytes_total counter
node_network_receive_bytes_total{device="docker0"} 3.2113968e+07
node_network_receive_bytes_total{device="enp0s3"} 5.75404803e+08
node_network_receive_bytes_total{device="lo"} 4.96457601e+08
node_network_receive_bytes_total{device="veth02b5642"} 8.972543e+06
node_network_receive_bytes_total{device="veth2466b1a"} 1.159047e+06
node_network_receive_bytes_total{device="veth40c6bad"} 4.343219e+06
node_network_receive_bytes_total{device="veth6999e49"} 2.581964e+06
node_network_receive_bytes_total{device="veth821ba89"} 7.5211803e+07
node_network_receive_bytes_total{device="vetha68ccf1"} 1.1204031e+07
node_network_receive_bytes_total{device="vetha9bee2f"} 3.957762e+06
```



# Metric types

Metric Type	Description
Counter	Cumulative metric representing a single increasing value
Gauges	Single numerical value which can either increase or decrease
Histograms	Samples values and counts them in configurable buckets
Summaries	Samples values and also provides sum and total counts





# How Prometheus is intended to be used

107

- Can be used to monitor most anything and analyze performance
  - Databases, VMs, servers, apps
- Works well for collecting and evaluating metrics
  - Numeric time series
- Good for dynamic service-oriented architectures and machine-centric monitoring
- Supports multi-dimensional data collection and querying
- Designed for reliability
- Each Prometheus server is standalone
  - Not dependent on remote services or network storage



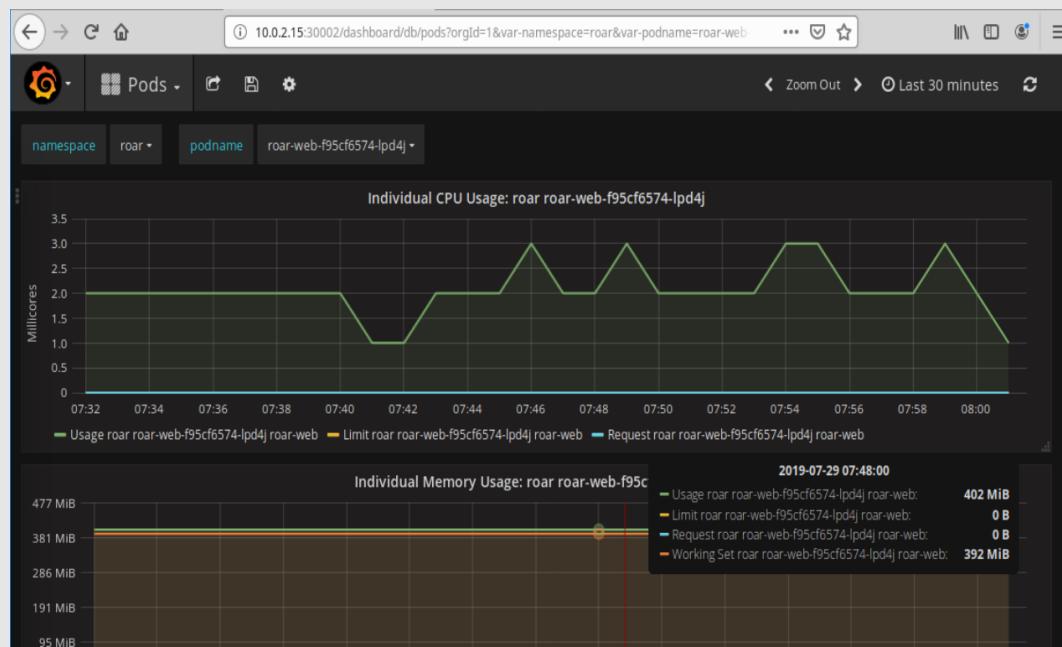
# Prometheus query examples

- count of pods per cluster and namespace
  - sum by (namespace) (kube\_pod\_info)
- number of containers by cluster and namespace without CPU limits (detects containers w/ no cpu limits)
  - count by (namespace)(sum by (namespace,pod,container)(kube\_pod\_container\_info{container!=""})) unless sum by (namespace,pod,container)(kube\_pod\_container\_resource\_limits{resource="cpu"})
- pod restarts by namespace
  - sum by (namespace)(changes(kube\_pod\_status\_ready{condition="true"}[5m]))
- pods not ready
  - sum by (namespace) (kube\_pod\_status\_ready{condition="false"})
- cpu overcommit
  - sum(kube\_pod\_container\_resource\_limits{resource="cpu"}) - sum(kube\_node\_status\_capacity\_cpu\_cores)
- memory overcommit
  - sum(kube\_pod\_container\_resource\_limits{resource="memory"}) - sum(kube\_node\_status\_capacity\_memory\_bytes)
- number of ready nodes per cluster
  - sum(kube\_node\_status\_condition{condition="Ready", status="true"}==1)
- nodes readiness flapping
  - sum(changes(kube\_node\_status\_condition{status="true"}, condition="Ready")[15m])) by (node) > 2
- cpu idle by cluster
  - sum((rate(container\_cpu\_usage\_seconds\_total{container!="POD", container!=""}[30m]) - on (namespace,pod,container) group\_left avg by (namespace,pod,container)(kube\_pod\_container\_resource\_requests{resource="cpu"})) \* -1 >0)
- memory idle by cluster
  - sum((container\_memory\_usage\_bytes{container!="POD", container!=""} - on (namespace,pod,container) avg by (namespace,pod,container)(kube\_pod\_container\_resource\_requests{resource="memory"})) \* -1 >0 ) / (1024\*1024\*1024)

- abs()
- absent()
- absent\_over\_time()
- ceil()
- changes()
- clamp()
- clamp\_max()
- clamp\_min()
- day\_of\_month()
- day\_of\_week()
- days\_in\_month()
- delta()
- deriv()
- exp()
- floor()
- histogram\_quantile()
- holt\_winters()
- hour()
- idelta()
- increase()
- rate()
- label\_join()
- label\_replace()
- ln()
- log2()
- log10()
- minute()
- month()
- predict\_linear()
- round()
- scalar()
- sgn()
- sort()
- sort\_desc()
- sqrt()
- time()
- timestamp()
- vector()
- year()
- <aggregation>\_over\_time()
- Trigonometric Functions



- Metric analytics and visualization suite
- Way to visualize time series for infrastructure and application analytics
- Allows users to build dashboards with graphs, panels, etc.
- Accepts data from many data sources including:
  - Graphite
  - Elasticsearch
  - InfluxDB
  - Prometheus
  - MySQL
  - Postgres
  - Cloudwatch





# Adding a Data Source

The screenshot displays the Grafana Configuration interface on the left and the Data Sources / Prometheus configuration page on the right.

**Left Panel (Configuration):**

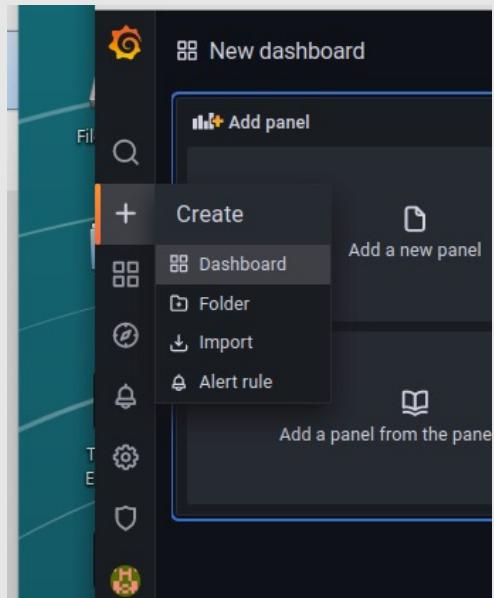
- Organization: Main Org.
- Navigation tabs: Data sources (highlighted), Users, Teams.
- Sidebar menu:
  - Configuration
  - Data sources (highlighted)
  - Users
  - Teams
  - Plugins
  - Preferences
- Exemplars section with a + Add button.
- Buttons at the bottom: Back, Explore, Delete, Save & test.

**Right Panel (Data Sources / Prometheus):**

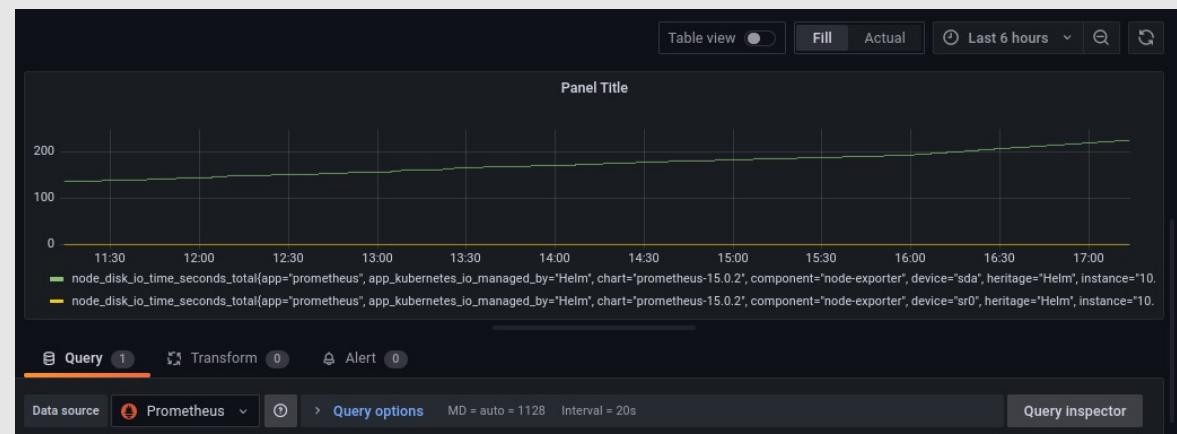
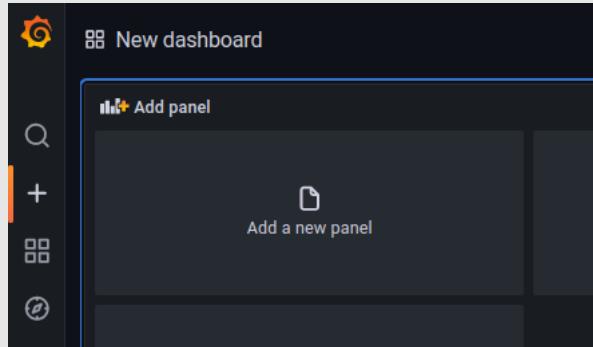
- Header: Data sources (highlighted), Users, Teams, Plugins, Preferences, API keys.
- Middle section:
  - No data sources defined.
  - Add data source button.
  - ProTip: You can also define data sources through configuration files. [Learn more](#).
- Details for the Prometheus data source:
  - Type: Prometheus
  - Settings tab (highlighted) and Dashboards tab.
  - Name: Prometheus (with a help icon).
  - Default toggle switch (turned on).
- HTTP section:
  - URL: rometheus-server.monitoring.svc.cluster.local:80
  - Access: Server (default)
  - Allowed cookies: New tag (enter key to add)
- Status message: Data source is working (green checkmark).
- Bottom buttons: Back, Explore, Delete, Save & test.



# Creating a Dashboard



This screenshot shows the Grafana query editor for a 'Time series' panel. At the top, there are buttons for 'Discard', 'Save', and 'Apply'. Below that is a search bar labeled 'Search options'. The main area is titled 'Query 1' and shows a single query: 'Metrics browser > node\_disk\_io\_time\_seconds\_total'. The data source is set to 'Prometheus'. The 'Query options' section indicates 'MD = auto = 1128' and 'Interval = 20s'. The chart area displays a green line graph titled 'Panel Title' showing disk I/O time over time, with two data series: 'node\_disk\_io\_time\_seconds\_total{app="prometheus", app\_kubernetes\_io\_managed\_by="Helm", chart="prometheus-15.0.2", component="node-exporter", device="sda", heritage="Helm", instance="10.' and 'node\_disk\_io\_time\_seconds\_total{app="prometheus", app\_kubernetes\_io\_managed\_by="Helm", chart="prometheus-15.0.2", component="node-exporter", device="sr0", heritage="Helm", instance="10.'.



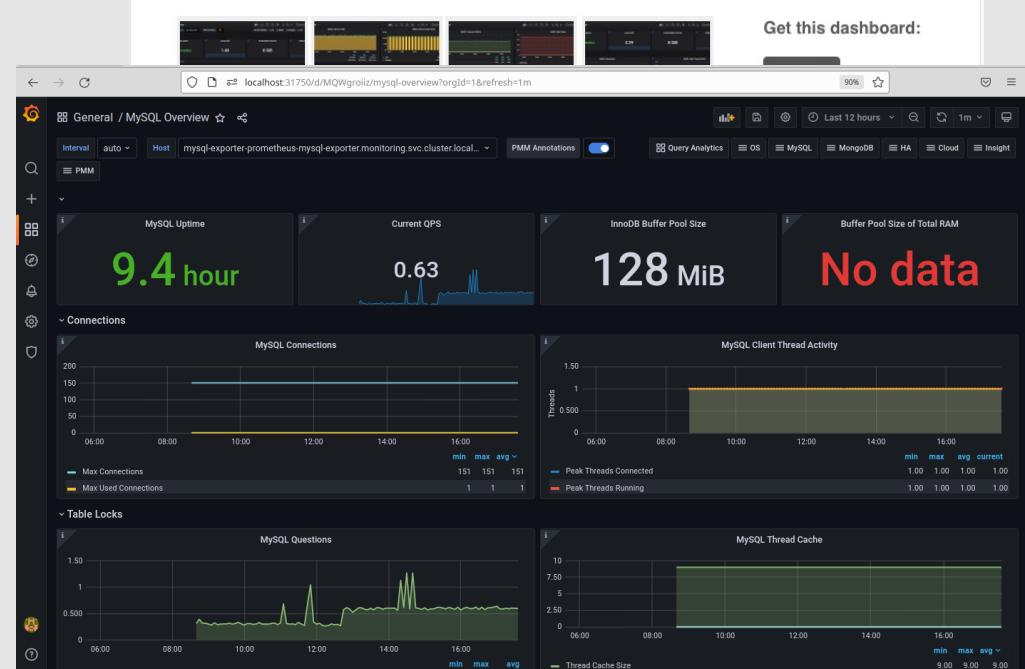


# Importing a Dashboard

The screenshot shows the 'Import' screen of the Grafana dashboard import interface. It includes a sidebar with icons for search, add, and dashboard management. The main area has two sections: 'Upload JSON file' with a blue 'Upload' button, and 'Import via grafana.com' with a URL input field containing 'https://grafana.com/grafana/dashboards/7362' and a 'Load' button.

The screenshot shows the 'MySQL Overview' dashboard by nasskach on the Grafana Labs website. It displays a summary card with 'Downloads: 69897' and 'Reviews: 9'. Below it is a detailed description: 'Dashboard from Percona Monitoring and Management project. Last updated: 3 years ago'. A 'Get this dashboard' button is visible at the bottom right.

The screenshot shows the 'Importing dashboard from Grafana.com' step of the import process. It displays metadata: 'Published by nasskach' and 'Updated on 2018-08-07 05:26'. Under 'Options', there are fields for 'Name' (set to 'MySQL Overview'), 'Folder' (set to 'General'), and 'Unique identifier (UID)'. The 'UID' section explains that it allows unique identification between multiple Grafana installs. At the bottom are 'Import' and 'Cancel' buttons.





# Additional Dashboards Available

- Node Exporter Full

<https://grafana.com/grafana/dashboards/1860>

Getting Started

Grafana Labs

All dashboards > Node Exporter Full

**Node Exporter Full** by rfraile

Downloads: 110573 Downloads: 110573

Reviews: 52 Reviews: 52

Last updated: 24 days ago Last updated: 24 days ago

Start with Grafana Cloud and the new FREE tier. Includes 10K series Prometheus or Graphite Metrics and 50gb Loki Logs

Add your review!

Overview Revisions Reviews

Get this dashboard:

1860 Copy ID to Clipboard

Nearly all default values exported by Prometheus node exporter graphed.

Only requires the default job\_name: node, add as many targets as you need in /etc/prometheus/prometheus.yml.

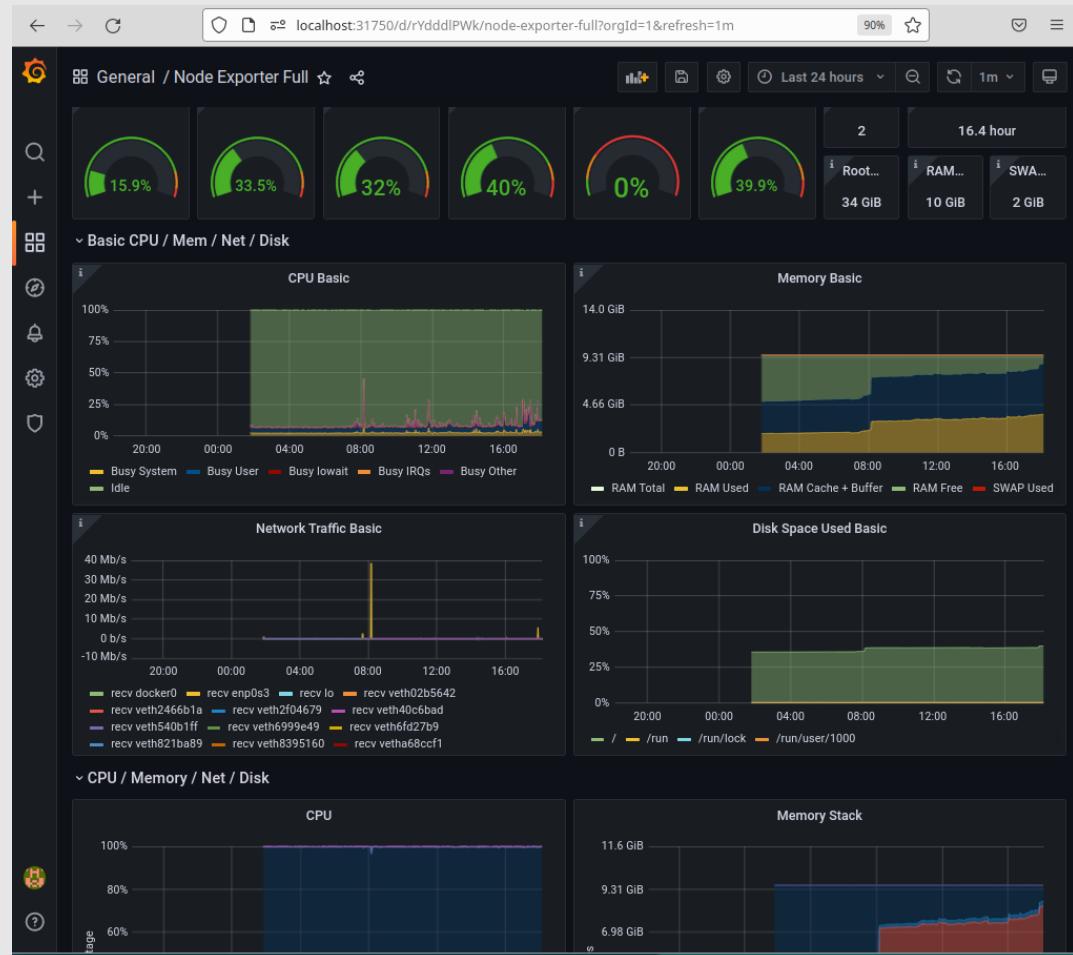
```
- job_name: node
  static_configs:
    - targets: ['localhost:9100']
```

Recommended for prometheus-node-exporter the arguments '--collector.systemd --collector.processes' because the graph uses some of their metrics.

Since revision 16, for prometheus-node-exporter v0.18 or newer. Since revision 12, for prometheus-node-exporter v0.16 or newer.

Dependencies:

- GRAFANA 7.3.7
- GAUGE
- GRAPH
- PROMETHEUS 1.0.0



## Lab 9 - Monitoring

Purpose: This lab will introduce you to a few of the ways we can monitor what is happening in our Kubernetes cluster and objects.

# RBAC



- Role-based Access Control

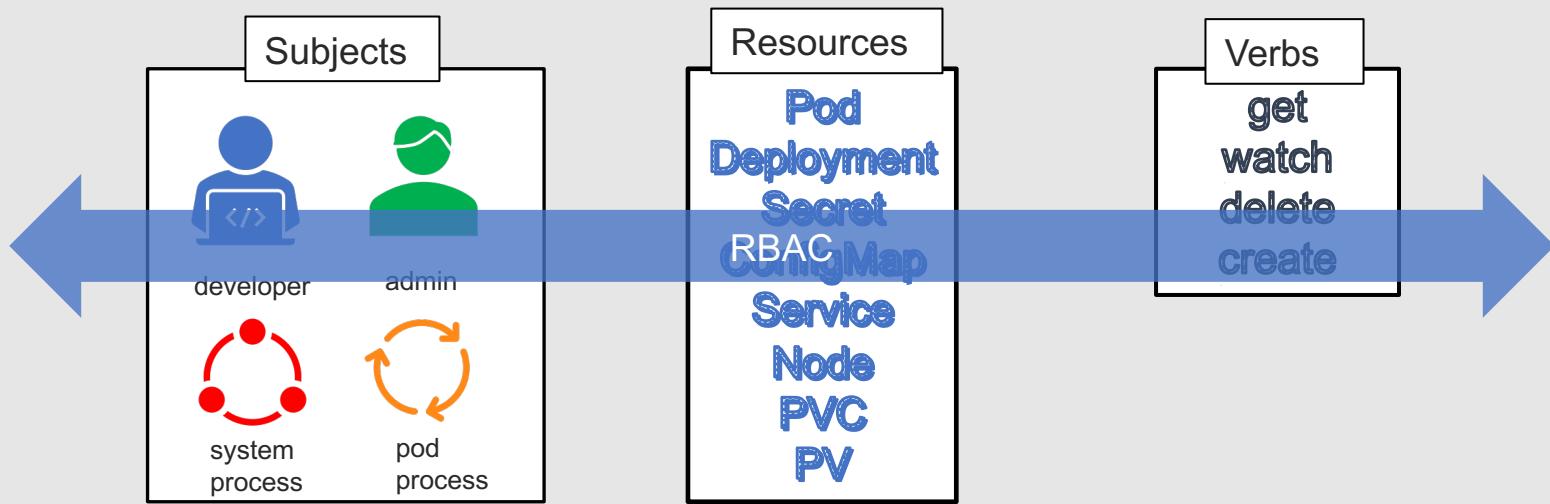
- Why do we need RBAC ? Because in a real-world k8s environment, we may need to:
  - Have multiple users with different properties, establishing a proper authentication mechanism.
  - Have full control over which operations each user or group of users can execute.
  - Have full control over which operations each process inside a pod can execute.
  - Limit the visibility of certain resources of namespaces.



- The pieces

- 3 types of elements involved

- Subjects: users or processes that need to access the API
- Resources: Kubernetes API Objects available in the cluster. (Examples include Pods, Deployments, Services, Nodes, and Persistent Volumes, etc.)
- Verbs: The Kubernetes operations that can be executed on the resources above. Multiple kinds of verbs available – all are CRUD operations (Examples include get, watch, create, delete, etc.)



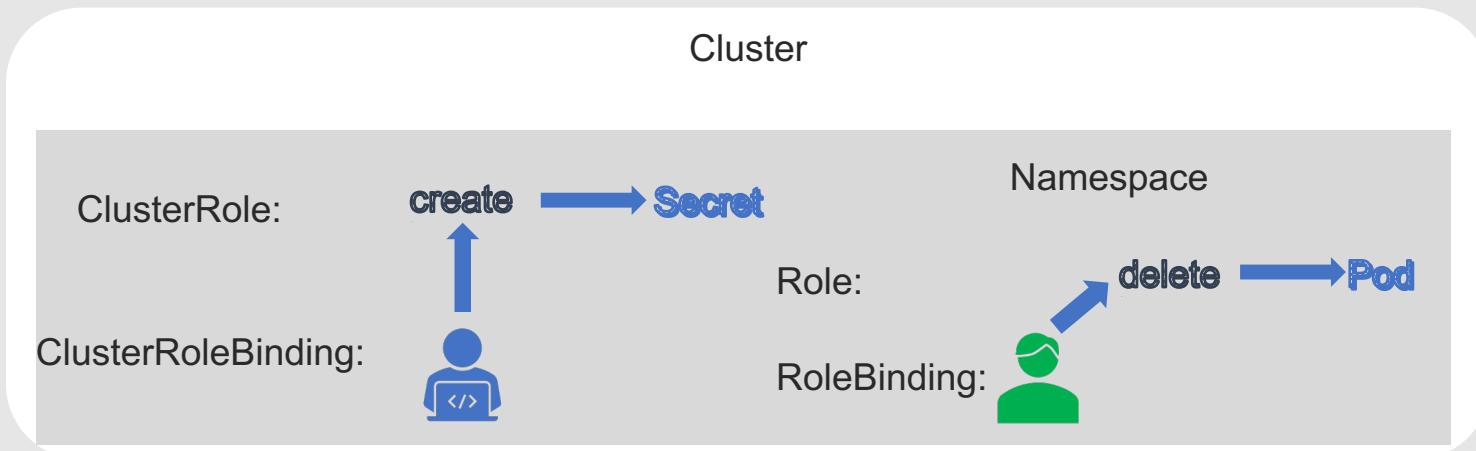
**Goal:** Connect subjects, API resources, and verbs. Specify, given a user, which operations can be executed over which resources.



- The mechanisms

- Object types

- k8s Roles: Connect API Resources and Verbs; can be reused for different Subjects; bound to a namespace
- k8s RoleBinding: Connects Subjects to Roles; Given a role that combines API objects and verbs, defines which subjects can use it.
- k8s ClusterRole: a role to be applied across a cluster
- k8s ClusterRoleBinding: Connects Subjects to ClusterRole.





# Kubernetes RBAC

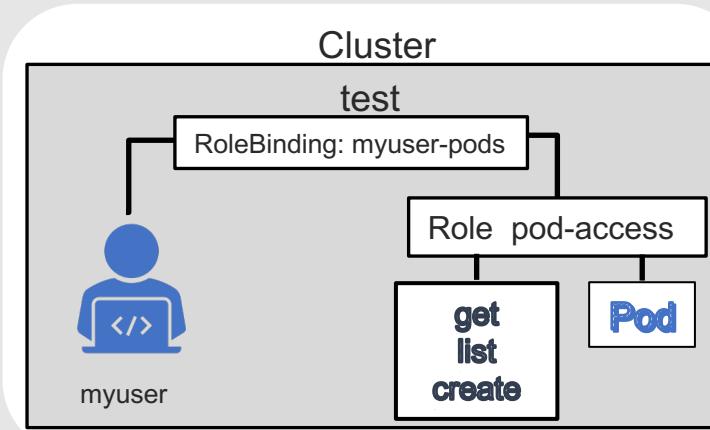
- Example: give user limited access to work with pods in test ns

- Specs define role and rolebindings

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: pod-access
  namespace: test
rules:
- apiGroups:
  - ""
resources:
- pods
verbs:
- get
- list
- create
```

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: myuser-pods
  namespace: test
subjects:
- kind: User
  name: myuser
roleRef:
  kind: Role
  name: pod-access
  apiGroup: rbac.authorization.k8s.io
```



- User can do these operations:
  - kubectl get pods –n test
  - kubectl describe pod –n test <name of pod>
  - kubectl create –f <pod-spec.yaml> -n test
- But not these
  - kubectl get pods –n default ~~(crossed out)~~
  - kubectl get pods –w –n test ~~(crossed out)~~



# Kubernetes RBAC - Service Accounts

120

- Users: These are global, and meant for humans or processes living outside the cluster.
- ServiceAccounts: These are namespaced and meant for intra-cluster processes running inside pods.
- Both authenticate against the API to get access to resources
- But there is no kubernetes “User” object
- So you can do “kubectl create serviceaccount <name>” but not “kubectl create user <name>”
- Bottom line: cluster does not store any information about users, so must be managed outside of cluster with certificates, tokens, etc.

## Create ServiceAccount

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: myapp
  labels:
    app: myapp
```

## Create Role

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: myapp-access-ep
  labels:
    app: myapp
rules:
- apiGroups: [""]
  resources: ["endpoints"]
  verbs: ["get"]
```

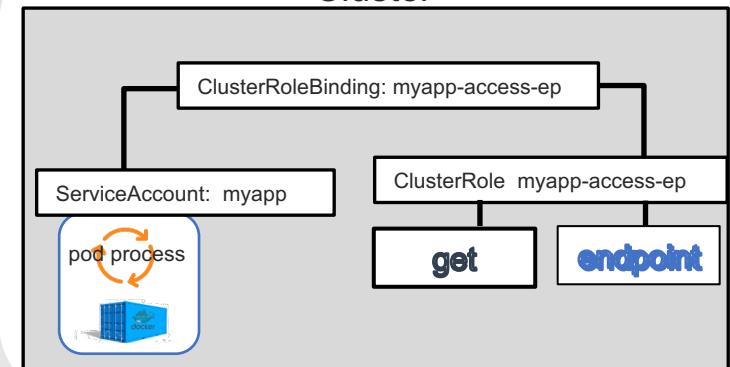
## Add ServiceAccount to Pod

```
[...]
spec:
  template:
    metadata:
      name: myapp
      labels:
        app: myapp
    spec:
      serviceAccountName: myapp
      containers:
[...]
```

## Tie Role to ServiceAccount

```
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: myapp-access-ep
  labels:
    app: myapp
subjects:
- kind: ServiceAccount
  name: myapp
roleRef:
  kind: ClusterRole
  name: myapp-access-ep
  apiGroup: rbac.authorization.k8s.io
```

## Cluster





# That's all - thanks!

The screenshot shows a web browser window with the URL [techskillstransformations.com](https://techskillstransformations.com). The page features a dark header with the company name "TECH SKILLS TRANSFORMATIONS, LLC". Below the header, a large banner on the left displays the title "TECH LEARNING MADE EASY" and a call-to-action button "Upcoming live training with O'Reilly Media!". To the right of the banner is a photograph of a person working at a desk with a laptop and a notebook. On the far right, there is a book cover for "Jenkins 2 Up & Running" by Brent Laster, featuring a fox illustration.

**Professional Git 1st Edition**  
by Brent Laster (Author)  
★★★★★ 7 customer reviews  
[Look inside](#)

**ABOUT US**

[techskillstransformations.com](https://techskillstransformations.com)  
[getskillsnow.com](https://getskillsnow.com)

Brent Laster